

In [1]:

```
"""
Author: Amruth Karun M V
Date: 19-Oct-2021
"""

import os
import pandas as pd
import numpy as np
import zipfile
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.models import Model
from tensorflow.keras.layers import (
    Input, Conv2D, MaxPool2D,
    AveragePooling2D, Flatten, GlobalAveragePooling2D,
    Dense, Dropout)
from keras.layers.merge import concatenate

from sklearn import metrics
import matplotlib.pyplot as plt
%matplotlib inline

TRAIN_PATH = "../input/covid19/"
EPOCHS = 100
BATCH_SIZE = 128
LEARNING_RATE = 0.001
INPUT_SIZE = (224, 224)

def load_data():
```

```
"""
```

```
Loads input data from directory
```

```
Arguments: None
```

```
Returns: Train and val generator
```

```
"""
```

```
train_datagen = keras.preprocessing.image.ImageDataGenerator(validation_split=0.2) # set validation split
```

```
train_generator = train_datagen.flow_from_directory(
```

```
    TRAIN_PATH,
```

```
    target_size=INPUT_SIZE,
```

```
    batch_size=BATCH_SIZE,
```

```
    shuffle=False,
```

```
    class_mode='categorical',
```

```
    subset='training') # set as training data
```

```
validation_generator = train_datagen.flow_from_directory(
```

```
    TRAIN_PATH,
```

```
    target_size=INPUT_SIZE,
```

```
    batch_size=BATCH_SIZE,
```

```
    shuffle=False,
```

```
    class_mode='categorical',
```

```
    subset='validation') # set as validation data
```

```
return train_generator, validation_generator
```

```
def inception_module(x, filters_1x1, filters_3x3_reduce, filters_3x3,  
                    filters_5x5_reduce, filters_5x5, filters_pool_proj,  
                    name=None):
```

```
"""
```

Represents an inception block

Arguments:

<i>x</i>	<i>-- input</i>
<i>filters_1x1</i>	<i>-- number of filters of the 1x1 convolutional layer in the first path</i>
<i>filters_3x3_reduce, filters_3x3</i>	<i>-- number of filters corresponding to the 1x1 and 3x3 convolutional layers in the second path</i>
<i>filters_5x5_reduce, filters_pool_proj</i>	<i>-- number of filters corresponding to the 1x1 and 5x5 convolutional layer in the third path</i>
<i>filters_pool_proj</i>	<i>-- number of filters of the 1x1 convolutional layer</i>

Returns: output layer

"""

```
conv_1x1 = Conv2D(filters_1x1, (1, 1), padding='same', activation='relu')(x)
```

```
conv_3x3 = Conv2D(filters_3x3_reduce, (1, 1), padding='same', activation='relu')(x)
```

```
conv_3x3 = Conv2D(filters_3x3, (3, 3), padding='same', activation='relu')(conv_3x3)
```

```
conv_5x5 = Conv2D(filters_5x5_reduce, (1, 1), padding='same', activation='relu')(x)
```

```
conv_5x5 = Conv2D(filters_5x5, (5, 5), padding='same', activation='relu')(conv_5x5)
```

```
pool_proj = MaxPool2D((3, 3), strides=(1, 1), padding='same')(x)
```

```
pool_proj = Conv2D(filters_pool_proj, (1, 1), padding='same', activation='relu')(pool_proj)
```

```
output = concatenate([conv_1x1, conv_3x3, conv_5x5, pool_proj], axis=3, name=name)
```

```
return output
```

```
def load_model():
```

"""

Creates a keras GoogleNet model

Arguments: None

Returns: GoogleNet Model

"""

```
input_layer = Input(shape=(224, 224, 3))
```

```
x = Conv2D(64, (7, 7), padding='same', strides=(2, 2), activation='relu', name='conv_1_7x7/2')(input_layer)
```

```
x = MaxPool2D((3, 3), padding='same', strides=(2, 2), name='max_pool_1_3x3/2')(x)
```

```
x = Conv2D(64, (1, 1), padding='same', strides=(1, 1), activation='relu', name='conv_2a_3x3/1')(x)
```

```
x = Conv2D(192, (3, 3), padding='same', strides=(1, 1), activation='relu', name='conv_2b_3x3/1')(x)
```

```
x = MaxPool2D((3, 3), padding='same', strides=(2, 2), name='max_pool_2_3x3/2')(x)
```

```
x = inception_module(x,
    filters_1x1=64,
    filters_3x3_reduce=96,
    filters_3x3=128,
    filters_5x5_reduce=16,
    filters_5x5=32,
    filters_pool_proj=32,
    name='inception_3a')
```

```
x = inception_module(x,
    filters_1x1=128,
    filters_3x3_reduce=128,
    filters_3x3=192,
    filters_5x5_reduce=32,
    filters_5x5=96,
    filters_pool_proj=64,
    name='inception_3b')
```

```
x = MaxPool2D((3, 3), padding='same', strides=(2, 2), name='max_pool_3_3x3/2')(x)
```

```
x = inception_module(x,  
    filters_1x1=192,  
    filters_3x3_reduce=96,  
    filters_3x3=208,  
    filters_5x5_reduce=16,  
    filters_5x5=48,  
    filters_pool_proj=64,  
    name='inception_4a')
```

```
x1 = AveragePooling2D((5, 5), strides=3)(x)  
x1 = Conv2D(128, (1, 1), padding='same', activation='relu')(x1)  
x1 = Flatten()(x1)  
x1 = Dense(1024, activation='relu')(x1)  
x1 = Dropout(0.7)(x1)  
x1 = Dense(3, activation='softmax', name='auxilliary_output_1')(x1)
```

```
x = inception_module(x,  
    filters_1x1=160,  
    filters_3x3_reduce=112,  
    filters_3x3=224,  
    filters_5x5_reduce=24,  
    filters_5x5=64,  
    filters_pool_proj=64,  
    name='inception_4b')
```

```
x = inception_module(x,  
    filters_1x1=128,  
    filters_3x3_reduce=128,
```

```
filters_3x3=256,  
filters_5x5_reduce=24,  
filters_5x5=64,  
filters_pool_proj=64,  
name='inception_4c')
```

```
x = inception_module(x,  
    filters_1x1=112,  
    filters_3x3_reduce=144,  
    filters_3x3=288,  
    filters_5x5_reduce=32,  
    filters_5x5=64,  
    filters_pool_proj=64,  
    name='inception_4d')
```

```
x2 = AveragePooling2D((5, 5), strides=3)(x)  
x2 = Conv2D(128, (1, 1), padding='same', activation='relu')(x2)  
x2 = Flatten()(x2)  
x2 = Dense(1024, activation='relu')(x2)  
x2 = Dropout(0.7)(x2)  
x2 = Dense(3, activation='softmax', name='auxilliary_output_2')(x2)
```

```
x = inception_module(x,  
    filters_1x1=256,  
    filters_3x3_reduce=160,  
    filters_3x3=320,  
    filters_5x5_reduce=32,  
    filters_5x5=128,  
    filters_pool_proj=128,  
    name='inception_4e')
```

```

x = MaxPool2D((3, 3), padding='same', strides=(2, 2), name='max_pool_4_3x3/2')(x)

x = inception_module(x,
                    filters_1x1=256,
                    filters_3x3_reduce=160,
                    filters_3x3=320,
                    filters_5x5_reduce=32,
                    filters_5x5=128,
                    filters_pool_proj=128,
                    name='inception_5a')

x = inception_module(x,
                    filters_1x1=384,
                    filters_3x3_reduce=192,
                    filters_3x3=384,
                    filters_5x5_reduce=48,
                    filters_5x5=128,
                    filters_pool_proj=128,
                    name='inception_5b')

x = GlobalAveragePooling2D(name='avg_pool_5_3x3/1')(x)

x = Dropout(0.4)(x)

x = Dense(3, activation='softmax', name='output')(x)
model = Model(input_layer, [x, x1, x2], name='GoogLeNet')
model.summary()

opt = Adam(learning_rate=LEARNING_RATE)
model.compile(loss = keras.losses.categorical_crossentropy, optimizer=opt, metrics=['accuracy'])

```

```
return model
```

```
def plot_curves(history):
```

```
    """
```

```
    Plots loss and accuracy and loss plots for  
    training and validation datasets
```

```
    Arguments:
```

```
        history -- training history
```

```
    Returns: None
```

```
    """
```

```
    plt.plot(history.history['loss'], color='b', label="Training loss")
```

```
    plt.plot(history.history['val_loss'], color='r', label="Validation loss")
```

```
    plt.legend()
```

```
    plt.title('Training Loss VS Validation Loss')
```

```
    plt.show()
```

```
    plt.plot(history.history['output_accuracy'], color='b', label="Training accuracy")
```

```
    plt.plot(history.history['val_output_accuracy'], color='r', label="Validation accuracy")
```

```
    plt.title('Training Accuracy VS Validation Accuracy')
```

```
    plt.legend()
```

```
    plt.show()
```

```
def get_confusion_matrix(model, data_generator):
```

```
    """
```

```
    Calculates the accuracy and displays the  
    confusion matrix for the input data
```

```
    Arguments:
```

```
        model          -- trained model
```


data_generator -- input data generator

Returns: None

"""

```
predictions = model.predict(data_generator, BATCH_SIZE)
```

```
y_pred = np.argmax(predictions[0], axis=1)
```

```
y_true = data_generator.classes
```

```
class_names = ['COVID', 'Normal', 'Pneumonia']
```

```
print("Score =", model.evaluate(data_generator, batch_size=BATCH_SIZE))
```

```
print("Accuracy = ", metrics.accuracy_score(y_true, y_pred))
```

```
cm = metrics.confusion_matrix(y_true, y_pred)
```

```
metrics.ConfusionMatrixDisplay(cm, display_labels=class_names).plot(cmap=plt.cm.Blues,  
                                                                    xticks_rotation='vertical')
```

```
plt.show()
```

```
def train_model(train_generator, val_generator):
```

"""

Trains GoogleNet model and saves the

trained weights to an H5 file.

Arguments:

train_generator -- train data generator

val_generator -- validation data generator

Returns: Trained model

"""

Loads the model

```
model = load_model()
```

```
earlystop = keras.callbacks.EarlyStopping(patience=10)
```

```
callbacks = [earlystop]
```

```
history = model.fit(  
    train_generator,  
    batch_size=BATCH_SIZE,  
    epochs=EPOCHS,  
    validation_data=val_generator,  
    validation_steps=val_generator.samples//BATCH_SIZE,  
    steps_per_epoch=train_generator.samples//BATCH_SIZE,  
    callbacks=callbacks)  
  
plot_curves(history)  
model.save_weights("model.h5")  
print("Model saved successfully!")  
  
return model
```

In [2]:

```
train_generator, val_generator = load_data()
model = train_model(train_generator, val_generator)

print("Confusion matrix for train data:")
get_confusion_matrix(model, train_generator)

print("Confusion matrix for val/test data:")
get_confusion_matrix(model, val_generator)
```

de zero

2021-10-19 09:48:41.497770: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1510] Created device /job:localhost/replica:0/task:0/device:GPU:0 with 15403 MB memory: -> device: 0, name: Tesla P100-PCIE-16GB, pci bus id: 0000:00:04.0, compute capability: 6.0

Model: "GoogLeNet"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	[(None, 224, 224, 3)]	0	

conv_1_7x7/2 (Conv2D)	(None, 112, 112, 64)	9472	input_1[0][0]

max_pool_1_3x3/2 (MaxPooling2D)	(None, 56, 56, 64)	0	conv_1_7x7/2[0][0]

conv_2a_3x3/1 (Conv2D)	(None, 56, 56, 64)	4160	max_pool_1_3x3/2[0][0]

conv_2b_3x3/1 (Conv2D)	(None, 56, 56, 192)	110784	conv_2a_3x3/1[0][0]

max_pool_2_3x3/2 (MaxPooling2D)	(None, 28, 28, 192)	0	conv_2b_3x3/1[0][0]

conv2d_1 (Conv2D)	(None, 28, 28, 96)	18528	max_pool_2_3x3/2[0][0]

conv2d_3 (Conv2D)	(None, 28, 28, 16)	3088	max_pool_2_3x3/2[0][0]

max_pooling2d (MaxPooling2D)	(None, 28, 28, 192)	0	max_pool_2_3x3/2[0][0]

conv2d (Conv2D)	(None, 28, 28, 64)	12352	max_pool_2_3x3/2[0][0]

conv2d_2 (Conv2D)	(None, 28, 28, 128)	110720	conv2d_1[0][0]

conv2d_4 (Conv2D)	(None, 28, 28, 32)	12832	conv2d_3[0][0]

conv2d_5 (Conv2D)	(None, 28, 28, 32)	6176	max_pooling2d[0][0]

inception_3a (Concatenate)	(None, 28, 28, 256)	0	conv2d[0][0] conv2d_2[0][0] conv2d_4[0][0] conv2d_5[0][0]
conv2d_7 (Conv2D)	(None, 28, 28, 128)	32896	inception_3a[0][0]
conv2d_9 (Conv2D)	(None, 28, 28, 32)	8224	inception_3a[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 28, 28, 256)	0	inception_3a[0][0]
conv2d_6 (Conv2D)	(None, 28, 28, 128)	32896	inception_3a[0][0]
conv2d_8 (Conv2D)	(None, 28, 28, 192)	221376	conv2d_7[0][0]
conv2d_10 (Conv2D)	(None, 28, 28, 96)	76896	conv2d_9[0][0]
conv2d_11 (Conv2D)	(None, 28, 28, 64)	16448	max_pooling2d_1[0][0]
inception_3b (Concatenate)	(None, 28, 28, 480)	0	conv2d_6[0][0] conv2d_8[0][0] conv2d_10[0][0] conv2d_11[0][0]
max_pool_3_3x3/2 (MaxPooling2D)	(None, 14, 14, 480)	0	inception_3b[0][0]
conv2d_13 (Conv2D)	(None, 14, 14, 96)	46176	max_pool_3_3x3/2[0][0]
conv2d_15 (Conv2D)	(None, 14, 14, 16)	7696	max_pool_3_3x3/2[0][0]
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 480)	0	max_pool_3_3x3/2[0][0]

conv2d_12 (Conv2D)	(None, 14, 14, 192)	92352	max_pool_3_3x3/2[0][0]
conv2d_14 (Conv2D)	(None, 14, 14, 208)	179920	conv2d_13[0][0]
conv2d_16 (Conv2D)	(None, 14, 14, 48)	19248	conv2d_15[0][0]
conv2d_17 (Conv2D)	(None, 14, 14, 64)	30784	max_pooling2d_2[0][0]
inception_4a (Concatenate)	(None, 14, 14, 512)	0	conv2d_12[0][0] conv2d_14[0][0] conv2d_16[0][0] conv2d_17[0][0]
conv2d_20 (Conv2D)	(None, 14, 14, 112)	57456	inception_4a[0][0]
conv2d_22 (Conv2D)	(None, 14, 14, 24)	12312	inception_4a[0][0]
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 512)	0	inception_4a[0][0]
conv2d_19 (Conv2D)	(None, 14, 14, 160)	82080	inception_4a[0][0]
conv2d_21 (Conv2D)	(None, 14, 14, 224)	226016	conv2d_20[0][0]
conv2d_23 (Conv2D)	(None, 14, 14, 64)	38464	conv2d_22[0][0]
conv2d_24 (Conv2D)	(None, 14, 14, 64)	32832	max_pooling2d_3[0][0]
inception_4b (Concatenate)	(None, 14, 14, 512)	0	conv2d_19[0][0] conv2d_21[0][0] conv2d_23[0][0]

			conv2d_24[0][0]
conv2d_26 (Conv2D)	(None, 14, 14, 128)	65664	inception_4b[0][0]
conv2d_28 (Conv2D)	(None, 14, 14, 24)	12312	inception_4b[0][0]
max_pooling2d_4 (MaxPooling2D)	(None, 14, 14, 512)	0	inception_4b[0][0]
conv2d_25 (Conv2D)	(None, 14, 14, 128)	65664	inception_4b[0][0]
conv2d_27 (Conv2D)	(None, 14, 14, 256)	295168	conv2d_26[0][0]
conv2d_29 (Conv2D)	(None, 14, 14, 64)	38464	conv2d_28[0][0]
conv2d_30 (Conv2D)	(None, 14, 14, 64)	32832	max_pooling2d_4[0][0]
inception_4c (Concatenate)	(None, 14, 14, 512)	0	conv2d_25[0][0] conv2d_27[0][0] conv2d_29[0][0] conv2d_30[0][0]
conv2d_32 (Conv2D)	(None, 14, 14, 144)	73872	inception_4c[0][0]
conv2d_34 (Conv2D)	(None, 14, 14, 32)	16416	inception_4c[0][0]
max_pooling2d_5 (MaxPooling2D)	(None, 14, 14, 512)	0	inception_4c[0][0]
conv2d_31 (Conv2D)	(None, 14, 14, 112)	57456	inception_4c[0][0]
conv2d_33 (Conv2D)	(None, 14, 14, 288)	373536	conv2d_32[0][0]

conv2d_35 (Conv2D)	(None, 14, 14, 64)	51264	conv2d_34[0][0]
conv2d_36 (Conv2D)	(None, 14, 14, 64)	32832	max_pooling2d_5[0][0]
inception_4d (Concatenate)	(None, 14, 14, 528)	0	conv2d_31[0][0] conv2d_33[0][0] conv2d_35[0][0] conv2d_36[0][0]
conv2d_39 (Conv2D)	(None, 14, 14, 160)	84640	inception_4d[0][0]
conv2d_41 (Conv2D)	(None, 14, 14, 32)	16928	inception_4d[0][0]
max_pooling2d_6 (MaxPooling2D)	(None, 14, 14, 528)	0	inception_4d[0][0]
conv2d_38 (Conv2D)	(None, 14, 14, 256)	135424	inception_4d[0][0]
conv2d_40 (Conv2D)	(None, 14, 14, 320)	461120	conv2d_39[0][0]
conv2d_42 (Conv2D)	(None, 14, 14, 128)	102528	conv2d_41[0][0]
conv2d_43 (Conv2D)	(None, 14, 14, 128)	67712	max_pooling2d_6[0][0]
inception_4e (Concatenate)	(None, 14, 14, 832)	0	conv2d_38[0][0] conv2d_40[0][0] conv2d_42[0][0] conv2d_43[0][0]
max_pool_4_3x3/2 (MaxPooling2D)	(None, 7, 7, 832)	0	inception_4e[0][0]
conv2d_45 (Conv2D)	(None, 7, 7, 160)	133280	max_pool_4_3x3/2[0][0]

conv2d_47 (Conv2D)	(None, 7, 7, 32)	26656	max_pool_4_3x3/2[0][0]
max_pooling2d_7 (MaxPooling2D)	(None, 7, 7, 832)	0	max_pool_4_3x3/2[0][0]
conv2d_44 (Conv2D)	(None, 7, 7, 256)	213248	max_pool_4_3x3/2[0][0]
conv2d_46 (Conv2D)	(None, 7, 7, 320)	461120	conv2d_45[0][0]
conv2d_48 (Conv2D)	(None, 7, 7, 128)	102528	conv2d_47[0][0]
conv2d_49 (Conv2D)	(None, 7, 7, 128)	106624	max_pooling2d_7[0][0]
inception_5a (Concatenate)	(None, 7, 7, 832)	0	conv2d_44[0][0] conv2d_46[0][0] conv2d_48[0][0] conv2d_49[0][0]
conv2d_51 (Conv2D)	(None, 7, 7, 192)	159936	inception_5a[0][0]
conv2d_53 (Conv2D)	(None, 7, 7, 48)	39984	inception_5a[0][0]
max_pooling2d_8 (MaxPooling2D)	(None, 7, 7, 832)	0	inception_5a[0][0]
average_pooling2d (AveragePooli	(None, 4, 4, 512)	0	inception_4a[0][0]
average_pooling2d_1 (AveragePoo	(None, 4, 4, 528)	0	inception_4d[0][0]
conv2d_50 (Conv2D)	(None, 7, 7, 384)	319872	inception_5a[0][0]
conv2d_52 (Conv2D)	(None, 7, 7, 384)	663936	conv2d_51[0][0]

conv2d_54 (Conv2D)	(None, 7, 7, 128)	153728	conv2d_53[0][0]
conv2d_55 (Conv2D)	(None, 7, 7, 128)	106624	max_pooling2d_8[0][0]
conv2d_18 (Conv2D)	(None, 4, 4, 128)	65664	average_pooling2d[0][0]
conv2d_37 (Conv2D)	(None, 4, 4, 128)	67712	average_pooling2d_1[0][0]
inception_5b (Concatenate)	(None, 7, 7, 1024)	0	conv2d_50[0][0] conv2d_52[0][0] conv2d_54[0][0] conv2d_55[0][0]
flatten (Flatten)	(None, 2048)	0	conv2d_18[0][0]
flatten_1 (Flatten)	(None, 2048)	0	conv2d_37[0][0]
avg_pool_5_3x3/1 (GlobalAveragePooling2D)	(None, 1024)	0	inception_5b[0][0]
dense (Dense)	(None, 1024)	2098176	flatten[0][0]
dense_1 (Dense)	(None, 1024)	2098176	flatten_1[0][0]
dropout_2 (Dropout)	(None, 1024)	0	avg_pool_5_3x3/1[0][0]
dropout (Dropout)	(None, 1024)	0	dense[0][0]
dropout_1 (Dropout)	(None, 1024)	0	dense_1[0][0]
output (Dense)	(None, 3)	3075	dropout_2[0][0]

```
-----
auxilliary_output_1 (Dense)      (None, 3)      3075      dropout[0][0]
-----
auxilliary_output_2 (Dense)      (None, 3)      3075      dropout_1[0][0]
=====
Total params: 10,312,505
Trainable params: 10,312,505
Non-trainable params: 0
-----
```

2021-10-19 09:48:43.564332: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the ML IR Optimization Passes are enabled (registered 2)

Epoch 1/100

2021-10-19 09:48:48.911158: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369] Loaded cuDNN version 8005

94/94 [=====] - 146s 1s/step - loss: 10.4544 - output_loss: 2.7401 - auxilliary_output_1_loss: 3.6297 - auxilliary_output_2_loss: 4.0847 - output_accuracy: 0.5730 - auxilliary_output_1_accuracy: 0.5526 - auxilliary_output_2_accuracy: 0.6082 - val_loss: 2.9125 - val_output_loss: 0.9392 - val_auxilliary_output_1_loss: 0.9952 - val_auxilliary_output_2_loss: 0.9781 - val_output_accuracy: 0.2626 - val_auxilliary_output_1_accuracy: 0.2456 - val_auxilliary_output_2_accuracy: 0.2456

Epoch 2/100

94/94 [=====] - 66s 705ms/step - loss: 4.2064 - output_loss: 1.0540 - auxilliary_output_1_loss: 1.5898 - auxilliary_output_2_loss: 1.5626 - output_accuracy: 0.6350 - auxilliary_output_1_accuracy: 0.6073 - auxilliary_output_2_accuracy: 0.6292 - val_loss: 2.8147 - val_output_loss: 0.8589 - val_auxilliary_output_1_loss: 0.9932 - val_auxilliary_output_2_loss: 0.9626 - val_output_accuracy: 0.6923 - val_auxilliary_output_1_accuracy: 0.2456 - val_auxilliary_output_2_accuracy: 0.2456

Epoch 3/100

94/94 [=====] - 66s 707ms/step - loss: 2.8171 - output_loss: 0.8671 - auxilliary_output_1_loss: 1.0495 - auxilliary_output_2_loss: 0.9005 - output_accuracy: 0.6691 - auxilliary_output_1_accuracy: 0.6084 - auxilliary_output_2_accuracy: 0.5723 - val_loss: 3.0545 - val_output_loss: 0.7702 - val_auxilliary_output_1_loss: 1.4307 - val_auxilliary_output_2_loss: 0.8536 - val_output_accuracy: 0.6923 - val_auxilliary_output_1_accuracy: 0.6923 - val_auxilliary_output_2_accuracy: 0.6923

Epoch 4/100

94/94 [=====] - 67s 712ms/step - loss: 11.5128 - output_loss: 1.6595 - auxilliary_output_1_loss: 4.7425 - auxilliary_output_2_loss: 5.1109 - output_accuracy: 0.6346 - auxilliary_output_1_accuracy: 0.5710 - auxilliary_output_2_accuracy: 0.6093 - val_loss: 2.3050 - val_output_loss: 0.7713 - val_auxilliary_output_1_loss: 0.7298 - val_auxilliary_output_2_loss: 0.8039 - val_output_accuracy: 0.6923 - val_auxilliary_output_1_accuracy: 0.6923 - val_auxilliary_output_2_accuracy: 0.6923

Epoch 5/100

94/94 [=====] - 67s 710ms/step - loss: 6.4662 - output_loss: 1.3789 - auxilliary_output_1_loss: 2.2838 - auxilliary_output_2_loss: 2.8035 - output_accuracy: 0.6805 - auxilliary_output_1_accuracy: 0.6596 - auxilliary_output_2_accuracy: 0.6763 - val_loss: 2.5897 - val_output_loss: 0.8612 - val_auxilliary_output_1_loss: 0.8369 - val_auxilliary_output_2_loss: 0.8915 - val_output_accuracy: 0.6923 - val_auxilliary_output_1_accuracy: 0.6912 - val_auxilliary_output_2_accuracy: 0.6923

Epoch 6/100

94/94 [=====] - 67s 714ms/step - loss: 3.5878 - output_loss: 0.8899 - auxilliary_output_1_loss: 1.7354 - auxilliary_output_2_loss: 0.9625 - output_accuracy: 0.6775 - auxilliary_output_1_accuracy: 0.5918 - auxilliary_output_2_accuracy: 0.6629 - val_loss: 2.3038 - val_output_loss: 0.7728 - val_auxilliary_output_1_loss: 0.7460 - val_auxilliary_output_2_loss: 0.7850 - val_output_accuracy: 0.6923 - val_auxilliary_output_1_accuracy: 0.6923 - val_auxilliary_output_2_accuracy: 0.6923

Epoch 7/100

94/94 [=====] - 67s 713ms/step - loss: 2.3564 - output_loss: 0.8313 - auxilliary_output_1_loss: 0.7299 - auxilliary_output_2_loss: 0.7952 - output_accuracy: 0.6691 - auxilliary_output_1_accuracy: 0.6789 - auxilliary_output_2_accuracy: 0.6658 - val_loss: 2.0938 - val_output_loss: 0.7285 - val_auxilliary_output_1_loss: 0.6771 - val_auxilliary_output_2_loss: 0.6882 - val_output_accuracy: 0.6923 - val_auxilliary_output_1_accuracy: 0.6797 - val_auxilliary_output_2_accuracy: 0.6970

Epoch 8/100

94/94 [=====] - 69s 729ms/step - loss: 1.9594 - output_loss: 0.7121 - auxilliary_output_1_loss: 0.6096 - auxilliary_output_2_loss: 0.6376 - output_accuracy: 0.6962 - auxilliary_output_1_accuracy: 0.6894 - auxilliary_output_2_accuracy: 0.6937 - val_loss: 1.8732 - val_output_loss: 0.6575 - val_auxilliary_output_1_loss: 0.5925 - val_auxilliary_output_2_loss: 0.6233 - val_output_accuracy: 0.6923 - val_auxilliary_output_1_accuracy: 0.7255 - val_auxilliary_output_2_accuracy: 0.7096

Epoch 9/100

94/94 [=====] - 69s 734ms/step - loss: 1.7716 - output_loss: 0.6315 - auxilliary_output_1_loss: 0.5588 - auxilliary_output_2_loss: 0.5814 - output_accuracy: 0.6964 - auxilliary_output_1_accuracy: 0.7395 - auxilliary_output_2_accuracy: 0.7193 - val_loss: 1.7408 - val_output_loss: 0.5853 - val_auxilliary_output_1_loss: 0.5698 - val_auxilliary_output_2_loss: 0.5858 - val_output_accuracy: 0.7249 - val_auxilliary_output_1_accuracy: 0.7853 - val_auxilliary_output_2_accuracy: 0.7310

Epoch 10/100

94/94 [=====] - 71s 755ms/step - loss: 1.7185 - output_loss: 0.6030 - auxilliary_output_1_loss: 0.5485 - auxilliary_output_2_loss: 0.5670 - output_accuracy: 0.7214 - auxilliary_output_1_accuracy: 0.7543 - auxilliary_output_2_accuracy: 0.7278 - val_loss: 1.8768 - val_output_loss: 0.6606 - val_auxilliary_output_1_loss: 0.5941 - val_auxilliary_output_2_loss: 0.6221 - val_output_accuracy: 0.7364 - val_auxilliary_output_1_accuracy: 0.7745 - val_auxilliary_output_2_accuracy: 0.7480

Epoch 11/100

94/94 [=====] - 72s 771ms/step - loss: 1.6171 - output_loss: 0.5741 - auxilliary_

output_1_loss: 0.5176 - auxilliary_output_2_loss: 0.5253 - output_accuracy: 0.7162 - auxilliary_output_1_accuracy: 0.7742 - auxilliary_output_2_accuracy: 0.7636 - val_loss: 1.7763 - val_output_loss: 0.5977 - val_auxilliary_output_1_loss: 0.5852 - val_auxilliary_output_2_loss: 0.5934 - val_output_accuracy: 0.7385 - val_auxilliary_output_1_accuracy: 0.7626 - val_auxilliary_output_2_accuracy: 0.7401

Epoch 12/100

94/94 [=====] - 72s 767ms/step - loss: 1.5291 - output_loss: 0.5458 - auxilliary_output_1_loss: 0.4855 - auxilliary_output_2_loss: 0.4978 - output_accuracy: 0.7211 - auxilliary_output_1_accuracy: 0.7848 - auxilliary_output_2_accuracy: 0.7729 - val_loss: 1.6041 - val_output_loss: 0.5643 - val_auxilliary_output_1_loss: 0.5223 - val_auxilliary_output_2_loss: 0.5175 - val_output_accuracy: 0.7446 - val_auxilliary_output_1_accuracy: 0.8268 - val_auxilliary_output_2_accuracy: 0.8268

Epoch 13/100

94/94 [=====] - 73s 775ms/step - loss: 1.3763 - output_loss: 0.4946 - auxilliary_output_1_loss: 0.4333 - auxilliary_output_2_loss: 0.4485 - output_accuracy: 0.7771 - auxilliary_output_1_accuracy: 0.8294 - auxilliary_output_2_accuracy: 0.8195 - val_loss: 1.6276 - val_output_loss: 0.6236 - val_auxilliary_output_1_loss: 0.4946 - val_auxilliary_output_2_loss: 0.5094 - val_output_accuracy: 0.7680 - val_auxilliary_output_1_accuracy: 0.8516 - val_auxilliary_output_2_accuracy: 0.8325

Epoch 14/100

94/94 [=====] - 70s 742ms/step - loss: 1.2973 - output_loss: 0.4998 - auxilliary_output_1_loss: 0.3859 - auxilliary_output_2_loss: 0.4116 - output_accuracy: 0.7820 - auxilliary_output_1_accuracy: 0.8409 - auxilliary_output_2_accuracy: 0.8288 - val_loss: 1.3575 - val_output_loss: 0.5236 - val_auxilliary_output_1_loss: 0.4053 - val_auxilliary_output_2_loss: 0.4286 - val_output_accuracy: 0.7469 - val_auxilliary_output_1_accuracy: 0.8655 - val_auxilliary_output_2_accuracy: 0.8461

Epoch 15/100

94/94 [=====] - 68s 726ms/step - loss: 1.3400 - output_loss: 0.5016 - auxilliary_output_1_loss: 0.4178 - auxilliary_output_2_loss: 0.4206 - output_accuracy: 0.7811 - auxilliary_output_1_accuracy: 0.8229 - auxilliary_output_2_accuracy: 0.8192 - val_loss: 1.4256 - val_output_loss: 0.5284 - val_auxilliary_output_1_loss: 0.4367 - val_auxilliary_output_2_loss: 0.4604 - val_output_accuracy: 0.7863 - val_auxilliary_output_1_accuracy: 0.8400 - val_auxilliary_output_2_accuracy: 0.8125

Epoch 16/100

94/94 [=====] - 73s 770ms/step - loss: 1.3080 - output_loss: 0.4789 - auxilliary_output_1_loss: 0.4059 - auxilliary_output_2_loss: 0.4232 - output_accuracy: 0.7866 - auxilliary_output_1_a

ccuracy: 0.8301 - auxilliary_output_2_accuracy: 0.8190 - val_loss: 1.4450 - val_output_loss: 0.5480 - val_ auxilliary_output_1_loss: 0.4400 - val_auxilliary_output_2_loss: 0.4569 - val_output_accuracy: 0.7993 - va l_auxilliary_output_1_accuracy: 0.8332 - val_auxilliary_output_2_accuracy: 0.8207

Epoch 17/100

94/94 [=====] - 73s 775ms/step - loss: 1.1820 - output_loss: 0.4326 - auxilliary_ output_1_loss: 0.3670 - auxilliary_output_2_loss: 0.3825 - output_accuracy: 0.8143 - auxilliary_output_1_a ccuracy: 0.8522 - auxilliary_output_2_accuracy: 0.8431 - val_loss: 1.4444 - val_output_loss: 0.5240 - val_ auxilliary_output_1_loss: 0.4648 - val_auxilliary_output_2_loss: 0.4556 - val_output_accuracy: 0.7846 - va l_auxilliary_output_1_accuracy: 0.8091 - val_auxilliary_output_2_accuracy: 0.8200

Epoch 18/100

94/94 [=====] - 71s 754ms/step - loss: 1.1653 - output_loss: 0.4263 - auxilliary_ output_1_loss: 0.3650 - auxilliary_output_2_loss: 0.3740 - output_accuracy: 0.8228 - auxilliary_output_1_a ccuracy: 0.8524 - auxilliary_output_2_accuracy: 0.8483 - val_loss: 1.3411 - val_output_loss: 0.5086 - val_ auxilliary_output_1_loss: 0.4078 - val_auxilliary_output_2_loss: 0.4247 - val_output_accuracy: 0.8482 - va l_auxilliary_output_1_accuracy: 0.8679 - val_auxilliary_output_2_accuracy: 0.8590

Epoch 19/100

94/94 [=====] - 68s 719ms/step - loss: 1.1365 - output_loss: 0.4476 - auxilliary_ output_1_loss: 0.3344 - auxilliary_output_2_loss: 0.3545 - output_accuracy: 0.8321 - auxilliary_output_1_a ccuracy: 0.8662 - auxilliary_output_2_accuracy: 0.8606 - val_loss: 1.2713 - val_output_loss: 0.4421 - val_ auxilliary_output_1_loss: 0.4014 - val_auxilliary_output_2_loss: 0.4278 - val_output_accuracy: 0.8196 - va l_auxilliary_output_1_accuracy: 0.8529 - val_auxilliary_output_2_accuracy: 0.8322

Epoch 20/100

94/94 [=====] - 68s 723ms/step - loss: 1.1609 - output_loss: 0.4169 - auxilliary_ output_1_loss: 0.3574 - auxilliary_output_2_loss: 0.3866 - output_accuracy: 0.8386 - auxilliary_output_1_a ccuracy: 0.8622 - auxilliary_output_2_accuracy: 0.8535 - val_loss: 1.2725 - val_output_loss: 0.4531 - val_ auxilliary_output_1_loss: 0.4079 - val_auxilliary_output_2_loss: 0.4116 - val_output_accuracy: 0.8346 - va l_auxilliary_output_1_accuracy: 0.8594 - val_auxilliary_output_2_accuracy: 0.8505

Epoch 21/100

94/94 [=====] - 68s 725ms/step - loss: 0.9767 - output_loss: 0.3589 - auxilliary_ output_1_loss: 0.3015 - auxilliary_output_2_loss: 0.3163 - output_accuracy: 0.8595 - auxilliary_output_1_a ccuracy: 0.8820 - auxilliary_output_2_accuracy: 0.8779 - val_loss: 1.2807 - val_output_loss: 0.4669 - val_

auxilliary_output_1_loss: 0.4034 - val_auxilliary_output_2_loss: 0.4104 - val_output_accuracy: 0.8295 - val_auxilliary_output_1_accuracy: 0.8954 - val_auxilliary_output_2_accuracy: 0.8781

Epoch 22/100

94/94 [=====] - 68s 721ms/step - loss: 0.9016 - output_loss: 0.3349 - auxilliary_output_1_loss: 0.2746 - auxilliary_output_2_loss: 0.2921 - output_accuracy: 0.8696 - auxilliary_output_1_accuracy: 0.8922 - auxilliary_output_2_accuracy: 0.8885 - val_loss: 0.9594 - val_output_loss: 0.3227 - val_auxilliary_output_1_loss: 0.3181 - val_auxilliary_output_2_loss: 0.3186 - val_output_accuracy: 0.8947 - val_auxilliary_output_1_accuracy: 0.9049 - val_auxilliary_output_2_accuracy: 0.8984

Epoch 23/100

94/94 [=====] - 68s 719ms/step - loss: 0.8581 - output_loss: 0.3158 - auxilliary_output_1_loss: 0.2603 - auxilliary_output_2_loss: 0.2820 - output_accuracy: 0.8769 - auxilliary_output_1_accuracy: 0.9012 - auxilliary_output_2_accuracy: 0.8930 - val_loss: 1.2069 - val_output_loss: 0.3849 - val_auxilliary_output_1_loss: 0.4161 - val_auxilliary_output_2_loss: 0.4058 - val_output_accuracy: 0.8641 - val_auxilliary_output_1_accuracy: 0.8512 - val_auxilliary_output_2_accuracy: 0.8668

Epoch 24/100

94/94 [=====] - 68s 719ms/step - loss: 0.8547 - output_loss: 0.3082 - auxilliary_output_1_loss: 0.2682 - auxilliary_output_2_loss: 0.2783 - output_accuracy: 0.8786 - auxilliary_output_1_accuracy: 0.8968 - auxilliary_output_2_accuracy: 0.8914 - val_loss: 1.3040 - val_output_loss: 0.4884 - val_auxilliary_output_1_loss: 0.4003 - val_auxilliary_output_2_loss: 0.4153 - val_output_accuracy: 0.8815 - val_auxilliary_output_1_accuracy: 0.8906 - val_auxilliary_output_2_accuracy: 0.8961

Epoch 25/100

94/94 [=====] - 67s 714ms/step - loss: 0.7474 - output_loss: 0.2633 - auxilliary_output_1_loss: 0.2360 - auxilliary_output_2_loss: 0.2481 - output_accuracy: 0.9010 - auxilliary_output_1_accuracy: 0.9132 - auxilliary_output_2_accuracy: 0.9086 - val_loss: 3.6459 - val_output_loss: 1.9100 - val_auxilliary_output_1_loss: 0.8462 - val_auxilliary_output_2_loss: 0.8897 - val_output_accuracy: 0.7548 - val_auxilliary_output_1_accuracy: 0.7721 - val_auxilliary_output_2_accuracy: 0.7704

Epoch 26/100

94/94 [=====] - 67s 710ms/step - loss: 1.1348 - output_loss: 0.4410 - auxilliary_output_1_loss: 0.3344 - auxilliary_output_2_loss: 0.3594 - output_accuracy: 0.8393 - auxilliary_output_1_accuracy: 0.8707 - auxilliary_output_2_accuracy: 0.8614 - val_loss: 1.4021 - val_output_loss: 0.5212 - val_auxilliary_output_1_loss: 0.4201 - val_auxilliary_output_2_loss: 0.4607 - val_output_accuracy: 0.8244 - va

l_auxilliary_output_1_accuracy: 0.8913 - val_auxilliary_output_2_accuracy: 0.8804

Epoch 27/100

94/94 [=====] - 67s 710ms/step - loss: 0.9151 - output_loss: 0.3422 - auxilliary_output_1_loss: 0.2763 - auxilliary_output_2_loss: 0.2966 - output_accuracy: 0.8668 - auxilliary_output_1_accuracy: 0.8941 - auxilliary_output_2_accuracy: 0.8861 - val_loss: 1.5706 - val_output_loss: 0.5607 - val_auxilliary_output_1_loss: 0.5141 - val_auxilliary_output_2_loss: 0.4958 - val_output_accuracy: 0.7843 - val_l_auxilliary_output_1_accuracy: 0.8091 - val_auxilliary_output_2_accuracy: 0.8190

Epoch 28/100

94/94 [=====] - 67s 708ms/step - loss: 0.9238 - output_loss: 0.3296 - auxilliary_output_1_loss: 0.2894 - auxilliary_output_2_loss: 0.3048 - output_accuracy: 0.8686 - auxilliary_output_1_accuracy: 0.8841 - auxilliary_output_2_accuracy: 0.8762 - val_loss: 1.0947 - val_output_loss: 0.3808 - val_auxilliary_output_1_loss: 0.3568 - val_auxilliary_output_2_loss: 0.3571 - val_output_accuracy: 0.8675 - val_l_auxilliary_output_1_accuracy: 0.8835 - val_auxilliary_output_2_accuracy: 0.8862

Epoch 29/100

94/94 [=====] - 67s 711ms/step - loss: 0.8259 - output_loss: 0.2960 - auxilliary_output_1_loss: 0.2596 - auxilliary_output_2_loss: 0.2703 - output_accuracy: 0.8871 - auxilliary_output_1_accuracy: 0.8993 - auxilliary_output_2_accuracy: 0.8970 - val_loss: 1.4568 - val_output_loss: 0.5047 - val_auxilliary_output_1_loss: 0.4585 - val_auxilliary_output_2_loss: 0.4936 - val_output_accuracy: 0.8740 - val_l_auxilliary_output_1_accuracy: 0.8668 - val_auxilliary_output_2_accuracy: 0.8607

Epoch 30/100

94/94 [=====] - 67s 710ms/step - loss: 0.6984 - output_loss: 0.2500 - auxilliary_output_1_loss: 0.2152 - auxilliary_output_2_loss: 0.2331 - output_accuracy: 0.9039 - auxilliary_output_1_accuracy: 0.9187 - auxilliary_output_2_accuracy: 0.9117 - val_loss: 1.0828 - val_output_loss: 0.3679 - val_auxilliary_output_1_loss: 0.3604 - val_auxilliary_output_2_loss: 0.3544 - val_output_accuracy: 0.9069 - val_l_auxilliary_output_1_accuracy: 0.9073 - val_auxilliary_output_2_accuracy: 0.9113

Epoch 31/100

94/94 [=====] - 66s 705ms/step - loss: 0.6892 - output_loss: 0.2488 - auxilliary_output_1_loss: 0.2136 - auxilliary_output_2_loss: 0.2268 - output_accuracy: 0.9038 - auxilliary_output_1_accuracy: 0.9187 - auxilliary_output_2_accuracy: 0.9125 - val_loss: 0.9428 - val_output_loss: 0.3267 - val_auxilliary_output_1_loss: 0.3079 - val_auxilliary_output_2_loss: 0.3082 - val_output_accuracy: 0.9022 - val_l_auxilliary_output_1_accuracy: 0.9120 - val_auxilliary_output_2_accuracy: 0.9076

Epoch 32/100

94/94 [=====] - 67s 710ms/step - loss: 0.6543 - output_loss: 0.2318 - auxilliary_output_1_loss: 0.2091 - auxilliary_output_2_loss: 0.2134 - output_accuracy: 0.9094 - auxilliary_output_1_accuracy: 0.9166 - auxilliary_output_2_accuracy: 0.9148 - val_loss: 1.1860 - val_output_loss: 0.4169 - val_auxilliary_output_1_loss: 0.3772 - val_auxilliary_output_2_loss: 0.3919 - val_output_accuracy: 0.8899 - val_auxilliary_output_1_accuracy: 0.9005 - val_auxilliary_output_2_accuracy: 0.8967

Epoch 33/100

94/94 [=====] - 67s 711ms/step - loss: 0.6329 - output_loss: 0.2194 - auxilliary_output_1_loss: 0.1965 - auxilliary_output_2_loss: 0.2170 - output_accuracy: 0.9140 - auxilliary_output_1_accuracy: 0.9241 - auxilliary_output_2_accuracy: 0.9196 - val_loss: 0.8897 - val_output_loss: 0.3060 - val_auxilliary_output_1_loss: 0.2906 - val_auxilliary_output_2_loss: 0.2932 - val_output_accuracy: 0.9113 - val_auxilliary_output_1_accuracy: 0.9198 - val_auxilliary_output_2_accuracy: 0.9178

Epoch 34/100

94/94 [=====] - 68s 723ms/step - loss: 0.8270 - output_loss: 0.2902 - auxilliary_output_1_loss: 0.2682 - auxilliary_output_2_loss: 0.2685 - output_accuracy: 0.8920 - auxilliary_output_1_accuracy: 0.8988 - auxilliary_output_2_accuracy: 0.8975 - val_loss: 1.6724 - val_output_loss: 0.5224 - val_auxilliary_output_1_loss: 0.6542 - val_auxilliary_output_2_loss: 0.4958 - val_output_accuracy: 0.7721 - val_auxilliary_output_1_accuracy: 0.6416 - val_auxilliary_output_2_accuracy: 0.7857

Epoch 35/100

94/94 [=====] - 74s 783ms/step - loss: 1.2889 - output_loss: 0.4009 - auxilliary_output_1_loss: 0.4983 - auxilliary_output_2_loss: 0.3896 - output_accuracy: 0.8147 - auxilliary_output_1_accuracy: 0.8248 - auxilliary_output_2_accuracy: 0.8320 - val_loss: 1.0953 - val_output_loss: 0.3817 - val_auxilliary_output_1_loss: 0.3559 - val_auxilliary_output_2_loss: 0.3577 - val_output_accuracy: 0.8648 - val_auxilliary_output_1_accuracy: 0.8869 - val_auxilliary_output_2_accuracy: 0.8723

Epoch 36/100

94/94 [=====] - 76s 804ms/step - loss: 0.6913 - output_loss: 0.2439 - auxilliary_output_1_loss: 0.2199 - auxilliary_output_2_loss: 0.2275 - output_accuracy: 0.9014 - auxilliary_output_1_accuracy: 0.9152 - auxilliary_output_2_accuracy: 0.9136 - val_loss: 1.0381 - val_output_loss: 0.3595 - val_auxilliary_output_1_loss: 0.3257 - val_auxilliary_output_2_loss: 0.3529 - val_output_accuracy: 0.9052 - val_auxilliary_output_1_accuracy: 0.9086 - val_auxilliary_output_2_accuracy: 0.9049

Epoch 37/100

94/94 [=====] - 75s 799ms/step - loss: 0.5583 - output_loss: 0.1961 - auxilliary_output_1_loss: 0.1764 - auxilliary_output_2_loss: 0.1858 - output_accuracy: 0.9217 - auxilliary_output_1_accuracy: 0.9335 - auxilliary_output_2_accuracy: 0.9301 - val_loss: 1.1141 - val_output_loss: 0.3846 - val_auxilliary_output_1_loss: 0.3536 - val_auxilliary_output_2_loss: 0.3758 - val_output_accuracy: 0.8421 - val_auxilliary_output_1_accuracy: 0.9076 - val_auxilliary_output_2_accuracy: 0.8451

Epoch 38/100

94/94 [=====] - 70s 739ms/step - loss: 0.6315 - output_loss: 0.2245 - auxilliary_output_1_loss: 0.2032 - auxilliary_output_2_loss: 0.2038 - output_accuracy: 0.9135 - auxilliary_output_1_accuracy: 0.9256 - auxilliary_output_2_accuracy: 0.9219 - val_loss: 1.1620 - val_output_loss: 0.3888 - val_auxilliary_output_1_loss: 0.3907 - val_auxilliary_output_2_loss: 0.3825 - val_output_accuracy: 0.8770 - val_auxilliary_output_1_accuracy: 0.8825 - val_auxilliary_output_2_accuracy: 0.8638

Epoch 39/100

94/94 [=====] - 67s 710ms/step - loss: 0.8293 - output_loss: 0.2853 - auxilliary_output_1_loss: 0.2745 - auxilliary_output_2_loss: 0.2695 - output_accuracy: 0.8964 - auxilliary_output_1_accuracy: 0.8998 - auxilliary_output_2_accuracy: 0.8983 - val_loss: 1.2375 - val_output_loss: 0.4359 - val_auxilliary_output_1_loss: 0.4079 - val_auxilliary_output_2_loss: 0.3937 - val_output_accuracy: 0.8699 - val_auxilliary_output_1_accuracy: 0.8818 - val_auxilliary_output_2_accuracy: 0.8753

Epoch 40/100

94/94 [=====] - 67s 710ms/step - loss: 0.6668 - output_loss: 0.2406 - auxilliary_output_1_loss: 0.2053 - auxilliary_output_2_loss: 0.2210 - output_accuracy: 0.9114 - auxilliary_output_1_accuracy: 0.9236 - auxilliary_output_2_accuracy: 0.9214 - val_loss: 1.0799 - val_output_loss: 0.3690 - val_auxilliary_output_1_loss: 0.3494 - val_auxilliary_output_2_loss: 0.3616 - val_output_accuracy: 0.9073 - val_auxilliary_output_1_accuracy: 0.9103 - val_auxilliary_output_2_accuracy: 0.8978

Epoch 41/100

94/94 [=====] - 67s 712ms/step - loss: 0.5282 - output_loss: 0.1885 - auxilliary_output_1_loss: 0.1627 - auxilliary_output_2_loss: 0.1770 - output_accuracy: 0.9281 - auxilliary_output_1_accuracy: 0.9390 - auxilliary_output_2_accuracy: 0.9352 - val_loss: 1.1469 - val_output_loss: 0.3996 - val_auxilliary_output_1_loss: 0.3599 - val_auxilliary_output_2_loss: 0.3874 - val_output_accuracy: 0.9018 - val_auxilliary_output_1_accuracy: 0.9151 - val_auxilliary_output_2_accuracy: 0.9052


Epoch 42/100


94/94 [=====] - 67s 717ms/step - loss: 0.4387 - output_loss: 0.1632 - auxilliary_

output_1_loss: 0.1315 - auxilliary_output_2_loss: 0.1441 - output_accuracy: 0.9371 - auxilliary_output_1_accuracy: 0.9516 - auxilliary_output_2_accuracy: 0.9457 - val_loss: 0.9355 - val_output_loss: 0.3220 - val_auxilliary_output_1_loss: 0.2999 - val_auxilliary_output_2_loss: 0.3136 - val_output_accuracy: 0.9222 - val_auxilliary_output_1_accuracy: 0.9321 - val_auxilliary_output_2_accuracy: 0.9300

Epoch 43/100

94/94 [=====] - 78s 830ms/step - loss: 0.3792 - output_loss: 0.1400 - auxilliary_output_1_loss: 0.1138 - auxilliary_output_2_loss: 0.1254 - output_accuracy: 0.9484 - auxilliary_output_1_accuracy: 0.9576 - auxilliary_output_2_accuracy: 0.9537 - val_loss: 0.9593 - val_output_loss: 0.3252 - val_auxilliary_output_1_loss: 0.3111 - val_auxilliary_output_2_loss: 0.3230 - val_output_accuracy: 0.9297 - val_auxilliary_output_1_accuracy: 0.9293 - val_auxilliary_output_2_accuracy: 0.9293






Model saved successfully!

Confusion matrix for train data:

95/95 [=====] - 52s 544ms/step - loss: 0.2735 - output_loss: 0.1026 - auxilliary_output_1_loss: 0.0816 - auxilliary_output_2_loss: 0.0893 - output_accuracy: 0.9595 - auxilliary_output_1_accuracy: 0.9677 - auxilliary_output_2_accuracy: 0.9666

Score = [0.2734605371952057, 0.10262028872966766, 0.08156668394804001, 0.08927353471517563, 0.959498465061877, 0.9676647782325745, 0.9665924310684204]

Accuracy = 0.9594984739750887



Confusion matrix for val/test data:

24/24 [=====] - 13s 534ms/step - loss: 0.9628 - output_loss: 0.3269 - auxilliary_output_1_loss: 0.3118 - auxilliary_output_2_loss: 0.3242 - output_accuracy: 0.9267 - auxilliary_output_1_accuracy: 0.9277 - auxilliary_output_2_accuracy: 0.9277

Score = [0.9628267884254456, 0.3268653452396393, 0.31176432967185974, 0.3241971433162689, 0.9267326593399048, 0.9277227520942688, 0.9277227520942688]

Accuracy = 0.9267326732673268

