

In [1]:

```
"""
Author: Amruth Karun M V
Date: 20-Oct-2021
"""

import os
import pandas as pd
import numpy as np
import zipfile
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.models import Model
from tensorflow.keras.layers import (
    Input, Conv2D, MaxPooling2D, Activation,
    AveragePooling2D, Flatten, BatchNormalization,
    Dense, Dropout, ZeroPadding2D, Add)
from keras.layers.merge import concatenate

from sklearn import metrics
import matplotlib.pyplot as plt
%matplotlib inline

TRAIN_PATH = "../input/covid19/"
EPOCHS = 100
BATCH_SIZE = 64
LEARNING_RATE = 0.001
INPUT_SIZE = (224, 224)

def load_data():
```

```
"""
```

```
Loads input data from directory
```

```
Arguments: None
```

```
Returns: Train and val generator
```

```
"""
```

```
train_datagen = keras.preprocessing.image.ImageDataGenerator(validation_split=0.2) # set validation split
```

```
train_generator = train_datagen.flow_from_directory(  
    TRAIN_PATH,  
    target_size=INPUT_SIZE,  
    batch_size=BATCH_SIZE,  
    shuffle=False,  
    class_mode='categorical',  
    subset='training') # set as training data
```

```
validation_generator = train_datagen.flow_from_directory(  
    TRAIN_PATH,  
    target_size=INPUT_SIZE,  
    batch_size=BATCH_SIZE,  
    shuffle=False,  
    class_mode='categorical',  
    subset='validation') # set as validation data
```

```
return train_generator, validation_generator
```

```
def identity_block(X, f, filters, stage, block):
```

```
    """
```

```
The identity block is the block that has no conv layer at shortcut.
```

```
Arguments:
```

```

X      -- input tensor
f      -- kernel size of middle conv layer at main path
filters -- list of integers, the filters of 3 conv layer at main path
stage  -- integer, current stage label, used for generating layer names
block  -- 'a','b'..., current block label, used for generating layer names
Returns: Output tensor for the block.
"""

conv_name_base = 'res' + str(stage) + block + '_branch'
bn_name_base = 'bn' + str(stage) + block + '_branch'
f1, f2, f3 = filters

X_shortcut = X

X = Conv2D(filters=f1, kernel_size=(1, 1), strides=(1, 1), padding='valid', name=conv_name_base + '2a')(X)
X = BatchNormalization(axis=3, name=bn_name_base + '2a')(X)
X = Activation('relu')(X)

X = Conv2D(filters=f2, kernel_size=(f, f), strides=(1, 1), padding='same', name=conv_name_base + '2b')(X)
X = BatchNormalization(axis=3, name=bn_name_base + '2b')(X)
X = Activation('relu')(X)

X = Conv2D(filters=f3, kernel_size=(1, 1), strides=(1, 1), padding='valid', name=conv_name_base + '2c')(X)
X = BatchNormalization(axis=3, name=bn_name_base + '2c')(X)

X = Add()([X, X_shortcut])# Skip Connection
X = Activation('relu')(X)

return X

```

```

def convolutional_block(X, f, filters, stage, block, strides=(2,2)):
    """
    A block that has a conv layer at shortcut.
    Arguments:
        X          -- input tensor
        f          -- the kernel size of middle conv layer at main path
        filters    -- list of integers, the filters of 3 conv layer at main path
        stage     -- integer, current stage label, used for generating layer names
        block     -- 'a','b'..., current block label, used for generating layer names
        strides    -- strides for the first conv layer in the block.
    Returns: Output tensor for the block.
    """

    conv_name_base = 'res' + str(stage) + block + '_branch'
    bn_name_base = 'bn' + str(stage) + block + '_branch'

    f1, f2, f3 = filters

    X_shortcut = X

    X = Conv2D(filters=f1, kernel_size=(1, 1), strides=strides, padding='valid', name=conv_name_base + '2a')(X)
    X = BatchNormalization(axis=3, name=bn_name_base + '2a')(X)
    X = Activation('relu')(X)

    X = Conv2D(filters=f2, kernel_size=(f, f), strides=(1, 1), padding='same', name=conv_name_base + '2b')(X)
    X = BatchNormalization(axis=3, name=bn_name_base + '2b')(X)
    X = Activation('relu')(X)

    X = Conv2D(filters=f3, kernel_size=(1, 1), strides=(1, 1), padding='valid', name=conv_name_base + '2c')(X)
    X = BatchNormalization(axis=3, name=bn_name_base + '2c')(X)

```

```

X_shortcut = Conv2D(filters=f3, kernel_size=(1, 1), strides=strides, padding='valid', name=conv_name_base
+ '1')(X_shortcut)
X_shortcut = BatchNormalization(axis=3, name=bn_name_base + '1')(X_shortcut)

X = Add()([X, X_shortcut])
X = Activation('relu')(X)

return X

```

```

def load_model():

```

```

    """

```

```

    Creates a keras ResNet-50 model

```

```

    Arguments: None

```

```

    Returns: ResNet-50 Model

```

```

    """

```

```

input_layer = Input(shape=(224, 224, 3))

```

```

X = ZeroPadding2D((3, 3))(input_layer)

```

```

X = Conv2D(64, (7, 7), strides=(2, 2), name='conv1')(X)

```

```

X = BatchNormalization(axis=3, name='bn_conv1')(X)

```

```

X = Activation('relu')(X)

```

```

X = MaxPooling2D((3, 3), strides=(2, 2))(X)

```

```

X = convolutional_block(X, f=3, filters=[64, 64, 256], stage=2, block='a', strides=(1,1))

```

```

X = identity_block(X, 3, [64, 64, 256], stage=2, block='b')

```

```

X = identity_block(X, 3, [64, 64, 256], stage=2, block='c')

```

```

X = convolutional_block(X, f=3, filters=[128, 128, 512], stage=3, block='a', strides=(2,2))
X = identity_block(X, 3, [128, 128, 512], stage=3, block='b')
X = identity_block(X, 3, [128, 128, 512], stage=3, block='c')
X = identity_block(X, 3, [128, 128, 512], stage=3, block='d')

X = convolutional_block(X, f=3, filters=[256, 256, 1024], stage=4, block='a', strides=(2,2))
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='b')
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='c')
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='d')
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='e')
X = identity_block(X, 3, [256, 256, 1024], stage=4, block='f')

X = convolutional_block(X, f=3, filters=[512, 512, 2048], stage=5, block='a', strides=(2,2))
X = identity_block(X, 3, [512, 512, 2048], stage=5, block='b')
X = identity_block(X, 3, [512, 512, 2048], stage=5, block='c')
X = AveragePooling2D(pool_size=(2, 2), padding='same')(X)
X = Dropout(0.4)(X)

# Define fully connected layers and output
X = Flatten()(X)
X = Dense(units=512, activation="relu")(X)
X = Dense(units=256, activation="relu")(X)
X = Dense(units=3, activation="softmax")(X)

model = Model(inputs=input_layer, outputs=X, name='ResNet50')
model.summary()

opt = Adam(learning_rate=LEARNING_RATE)
model.compile(loss = keras.losses.categorical_crossentropy, optimizer=opt, metrics=['accuracy'])

```

```
return model
```

```
def plot_curves(history):
```

```
    """
```

```
    Plots loss and accuracy and loss plots for  
    training and validation datasets
```

```
    Arguments:
```

```
        history -- training history
```

```
    Returns: None
```

```
    """
```

```
    plt.plot(history.history['loss'], color='b', label="Training loss")
```

```
    plt.plot(history.history['val_loss'], color='r', label="Validation loss")
```

```
    plt.legend()
```

```
    plt.title('Training Loss VS Validation Loss')
```

```
    plt.show()
```

```
    plt.plot(history.history['accuracy'], color='b', label="Training accuracy")
```

```
    plt.plot(history.history['val_accuracy'], color='r', label="Validation accuracy")
```

```
    plt.title('Training Accuracy VS Validation Accuracy')
```

```
    plt.legend()
```

```
    plt.show()
```

```
def get_confusion_matrix(model, data_generator):
```

```
    """
```

```
    Calculates the accuracy and displays the  
    confusion matrix for the input data
```

```
    Arguments:
```

```
        model          -- trained model
```

```
data_generator -- input data generator
```

```
Returns: None
```

```
"""
```

```
predictions = model.predict(data_generator, BATCH_SIZE)
```

```
y_pred = np.argmax(predictions, axis=1)
```

```
y_true = data_generator.classes
```

```
class_names = ['COVID', 'Normal', 'Pneumonia']
```

```
print("Score =", model.evaluate(data_generator, batch_size=BATCH_SIZE))
```

```
print("Accuracy = ", metrics.accuracy_score(y_true, y_pred))
```

```
cm = metrics.confusion_matrix(y_true, y_pred)
```

```
metrics.ConfusionMatrixDisplay(cm, display_labels=class_names).plot(cmap=plt.cm.Blues,  
                                                                    xticks_rotation='vertical')
```

```
plt.show()
```

```
def train_model(train_generator, val_generator):
```

```
    """
```

```
Trains ResNet-50 model and saves the
```

```
trained weights to an H5 file.
```

```
Arguments:
```

```
train_generator -- train data generator
```

```
val_generator -- validation data generator
```

```
Returns: Trained model
```

```
    """
```

```
# Loads the model
```

```
model = load_model()
```

```
earlystop = keras.callbacks.EarlyStopping(monitor='val_accuracy', mode='max', verbose=1, patience=10)
```

```
callbacks = [earlystop]
```



```
history = model.fit(  
    train_generator,  
    batch_size=BATCH_SIZE,  
    epochs=EPOCHS,  
    validation_data=val_generator,  
    validation_steps=val_generator.samples//BATCH_SIZE,  
    steps_per_epoch=train_generator.samples//BATCH_SIZE,  
    callbacks=callbacks)  
  
plot_curves(history)  
model.save_weights("model_resnet50.h5")  
print("Model saved successfully!")  
  
return model
```

In [2]:

```
train_generator, val_generator = load_data()
model = train_model(train_generator, val_generator)

print("Confusion matrix for train data:")
get_confusion_matrix(model, train_generator)

print("Confusion matrix for val/test data:")
get_confusion_matrix(model, val_generator)
```

Model: "ResNet50"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	[(None, 224, 224, 3)]	0	

zero_padding2d (ZeroPadding2D)	(None, 230, 230, 3)	0	input_1[0][0]

conv1 (Conv2D)	(None, 112, 112, 64)	9472	zero_padding2d[0][0]

bn_conv1 (BatchNormalization)	(None, 112, 112, 64)	256	conv1[0][0]

activation (Activation)	(None, 112, 112, 64)	0	bn_conv1[0][0]

max_pooling2d (MaxPooling2D)	(None, 55, 55, 64)	0	activation[0][0]

res2a_branch2a (Conv2D)	(None, 55, 55, 64)	4160	max_pooling2d[0][0]

bn2a_branch2a (BatchNormalization)	(None, 55, 55, 64)	256	res2a_branch2a[0][0]

activation_1 (Activation)	(None, 55, 55, 64)	0	bn2a_branch2a[0][0]

res2a_branch2b (Conv2D)	(None, 55, 55, 64)	36928	activation_1[0][0]

bn2a_branch2b (BatchNormalization)	(None, 55, 55, 64)	256	res2a_branch2b[0][0]

activation_2 (Activation)	(None, 55, 55, 64)	0	bn2a_branch2b[0][0]

res2a_branch2c (Conv2D)	(None, 55, 55, 256)	16640	activation_2[0][0]

res2a_branch1 (Conv2D)	(None, 55, 55, 256)	16640	max_pooling2d[0][0]
bn2a_branch2c (BatchNormalizati	(None, 55, 55, 256)	1024	res2a_branch2c[0][0]
bn2a_branch1 (BatchNormalizatio	(None, 55, 55, 256)	1024	res2a_branch1[0][0]
add (Add)	(None, 55, 55, 256)	0	bn2a_branch2c[0][0] bn2a_branch1[0][0]
activation_3 (Activation)	(None, 55, 55, 256)	0	add[0][0]
res2b_branch2a (Conv2D)	(None, 55, 55, 64)	16448	activation_3[0][0]
bn2b_branch2a (BatchNormalizati	(None, 55, 55, 64)	256	res2b_branch2a[0][0]
activation_4 (Activation)	(None, 55, 55, 64)	0	bn2b_branch2a[0][0]
res2b_branch2b (Conv2D)	(None, 55, 55, 64)	36928	activation_4[0][0]
bn2b_branch2b (BatchNormalizati	(None, 55, 55, 64)	256	res2b_branch2b[0][0]
activation_5 (Activation)	(None, 55, 55, 64)	0	bn2b_branch2b[0][0]
res2b_branch2c (Conv2D)	(None, 55, 55, 256)	16640	activation_5[0][0]
bn2b_branch2c (BatchNormalizati	(None, 55, 55, 256)	1024	res2b_branch2c[0][0]
add_1 (Add)	(None, 55, 55, 256)	0	bn2b_branch2c[0][0] activation_3[0][0]
activation_6 (Activation)	(None, 55, 55, 256)	0	add_1[0][0]

res2c_branch2a (Conv2D)	(None, 55, 55, 64)	16448	activation_6[0][0]
bn2c_branch2a (BatchNormalizati	(None, 55, 55, 64)	256	res2c_branch2a[0][0]
activation_7 (Activation)	(None, 55, 55, 64)	0	bn2c_branch2a[0][0]
res2c_branch2b (Conv2D)	(None, 55, 55, 64)	36928	activation_7[0][0]
bn2c_branch2b (BatchNormalizati	(None, 55, 55, 64)	256	res2c_branch2b[0][0]
activation_8 (Activation)	(None, 55, 55, 64)	0	bn2c_branch2b[0][0]
res2c_branch2c (Conv2D)	(None, 55, 55, 256)	16640	activation_8[0][0]
bn2c_branch2c (BatchNormalizati	(None, 55, 55, 256)	1024	res2c_branch2c[0][0]
add_2 (Add)	(None, 55, 55, 256)	0	bn2c_branch2c[0][0] activation_6[0][0]
activation_9 (Activation)	(None, 55, 55, 256)	0	add_2[0][0]
res3a_branch2a (Conv2D)	(None, 28, 28, 128)	32896	activation_9[0][0]
bn3a_branch2a (BatchNormalizati	(None, 28, 28, 128)	512	res3a_branch2a[0][0]
activation_10 (Activation)	(None, 28, 28, 128)	0	bn3a_branch2a[0][0]
res3a_branch2b (Conv2D)	(None, 28, 28, 128)	147584	activation_10[0][0]
bn3a_branch2b (BatchNormalizati	(None, 28, 28, 128)	512	res3a_branch2b[0][0]

activation_11 (Activation)	(None, 28, 28, 128)	0	bn3a_branch2b[0][0]
res3a_branch2c (Conv2D)	(None, 28, 28, 512)	66048	activation_11[0][0]
res3a_branch1 (Conv2D)	(None, 28, 28, 512)	131584	activation_9[0][0]
bn3a_branch2c (BatchNormalizati	(None, 28, 28, 512)	2048	res3a_branch2c[0][0]
bn3a_branch1 (BatchNormalizatio	(None, 28, 28, 512)	2048	res3a_branch1[0][0]
add_3 (Add)	(None, 28, 28, 512)	0	bn3a_branch2c[0][0] bn3a_branch1[0][0]
activation_12 (Activation)	(None, 28, 28, 512)	0	add_3[0][0]
res3b_branch2a (Conv2D)	(None, 28, 28, 128)	65664	activation_12[0][0]
bn3b_branch2a (BatchNormalizati	(None, 28, 28, 128)	512	res3b_branch2a[0][0]
activation_13 (Activation)	(None, 28, 28, 128)	0	bn3b_branch2a[0][0]
res3b_branch2b (Conv2D)	(None, 28, 28, 128)	147584	activation_13[0][0]
bn3b_branch2b (BatchNormalizati	(None, 28, 28, 128)	512	res3b_branch2b[0][0]
activation_14 (Activation)	(None, 28, 28, 128)	0	bn3b_branch2b[0][0]
res3b_branch2c (Conv2D)	(None, 28, 28, 512)	66048	activation_14[0][0]
bn3b_branch2c (BatchNormalizati	(None, 28, 28, 512)	2048	res3b_branch2c[0][0]

add_4 (Add)	(None, 28, 28, 512)	0	bn3b_branch2c[0][0] activation_12[0][0]
activation_15 (Activation)	(None, 28, 28, 512)	0	add_4[0][0]
res3c_branch2a (Conv2D)	(None, 28, 28, 128)	65664	activation_15[0][0]
bn3c_branch2a (BatchNormalizati	(None, 28, 28, 128)	512	res3c_branch2a[0][0]
activation_16 (Activation)	(None, 28, 28, 128)	0	bn3c_branch2a[0][0]
res3c_branch2b (Conv2D)	(None, 28, 28, 128)	147584	activation_16[0][0]
bn3c_branch2b (BatchNormalizati	(None, 28, 28, 128)	512	res3c_branch2b[0][0]
activation_17 (Activation)	(None, 28, 28, 128)	0	bn3c_branch2b[0][0]
res3c_branch2c (Conv2D)	(None, 28, 28, 512)	66048	activation_17[0][0]
bn3c_branch2c (BatchNormalizati	(None, 28, 28, 512)	2048	res3c_branch2c[0][0]
add_5 (Add)	(None, 28, 28, 512)	0	bn3c_branch2c[0][0] activation_15[0][0]
activation_18 (Activation)	(None, 28, 28, 512)	0	add_5[0][0]
res3d_branch2a (Conv2D)	(None, 28, 28, 128)	65664	activation_18[0][0]
bn3d_branch2a (BatchNormalizati	(None, 28, 28, 128)	512	res3d_branch2a[0][0]

activation_19 (Activation)	(None, 28, 28, 128)	0	bn3d_branch2a[0][0]
res3d_branch2b (Conv2D)	(None, 28, 28, 128)	147584	activation_19[0][0]
bn3d_branch2b (BatchNormalizati	(None, 28, 28, 128)	512	res3d_branch2b[0][0]
activation_20 (Activation)	(None, 28, 28, 128)	0	bn3d_branch2b[0][0]
res3d_branch2c (Conv2D)	(None, 28, 28, 512)	66048	activation_20[0][0]
bn3d_branch2c (BatchNormalizati	(None, 28, 28, 512)	2048	res3d_branch2c[0][0]
add_6 (Add)	(None, 28, 28, 512)	0	bn3d_branch2c[0][0] activation_18[0][0]
activation_21 (Activation)	(None, 28, 28, 512)	0	add_6[0][0]
res4a_branch2a (Conv2D)	(None, 14, 14, 256)	131328	activation_21[0][0]
bn4a_branch2a (BatchNormalizati	(None, 14, 14, 256)	1024	res4a_branch2a[0][0]
activation_22 (Activation)	(None, 14, 14, 256)	0	bn4a_branch2a[0][0]
res4a_branch2b (Conv2D)	(None, 14, 14, 256)	590080	activation_22[0][0]
bn4a_branch2b (BatchNormalizati	(None, 14, 14, 256)	1024	res4a_branch2b[0][0]
activation_23 (Activation)	(None, 14, 14, 256)	0	bn4a_branch2b[0][0]
res4a_branch2c (Conv2D)	(None, 14, 14, 1024)	263168	activation_23[0][0]

res4a_branch1 (Conv2D)	(None, 14, 14, 1024)	525312	activation_21[0][0]
bn4a_branch2c (BatchNormalizati	(None, 14, 14, 1024)	4096	res4a_branch2c[0][0]
bn4a_branch1 (BatchNormalizatio	(None, 14, 14, 1024)	4096	res4a_branch1[0][0]
add_7 (Add)	(None, 14, 14, 1024)	0	bn4a_branch2c[0][0] bn4a_branch1[0][0]
activation_24 (Activation)	(None, 14, 14, 1024)	0	add_7[0][0]
res4b_branch2a (Conv2D)	(None, 14, 14, 256)	262400	activation_24[0][0]
bn4b_branch2a (BatchNormalizati	(None, 14, 14, 256)	1024	res4b_branch2a[0][0]
activation_25 (Activation)	(None, 14, 14, 256)	0	bn4b_branch2a[0][0]
res4b_branch2b (Conv2D)	(None, 14, 14, 256)	590080	activation_25[0][0]
bn4b_branch2b (BatchNormalizati	(None, 14, 14, 256)	1024	res4b_branch2b[0][0]
activation_26 (Activation)	(None, 14, 14, 256)	0	bn4b_branch2b[0][0]
res4b_branch2c (Conv2D)	(None, 14, 14, 1024)	263168	activation_26[0][0]
bn4b_branch2c (BatchNormalizati	(None, 14, 14, 1024)	4096	res4b_branch2c[0][0]
add_8 (Add)	(None, 14, 14, 1024)	0	bn4b_branch2c[0][0] activation_24[0][0]
activation_27 (Activation)	(None, 14, 14, 1024)	0	add_8[0][0]

res4c_branch2a (Conv2D)	(None, 14, 14, 256)	262400	activation_27[0][0]
bn4c_branch2a (BatchNormalizati	(None, 14, 14, 256)	1024	res4c_branch2a[0][0]
activation_28 (Activation)	(None, 14, 14, 256)	0	bn4c_branch2a[0][0]
res4c_branch2b (Conv2D)	(None, 14, 14, 256)	590080	activation_28[0][0]
bn4c_branch2b (BatchNormalizati	(None, 14, 14, 256)	1024	res4c_branch2b[0][0]
activation_29 (Activation)	(None, 14, 14, 256)	0	bn4c_branch2b[0][0]
res4c_branch2c (Conv2D)	(None, 14, 14, 1024)	263168	activation_29[0][0]
bn4c_branch2c (BatchNormalizati	(None, 14, 14, 1024)	4096	res4c_branch2c[0][0]
add_9 (Add)	(None, 14, 14, 1024)	0	bn4c_branch2c[0][0] activation_27[0][0]
activation_30 (Activation)	(None, 14, 14, 1024)	0	add_9[0][0]
res4d_branch2a (Conv2D)	(None, 14, 14, 256)	262400	activation_30[0][0]
bn4d_branch2a (BatchNormalizati	(None, 14, 14, 256)	1024	res4d_branch2a[0][0]
activation_31 (Activation)	(None, 14, 14, 256)	0	bn4d_branch2a[0][0]
res4d_branch2b (Conv2D)	(None, 14, 14, 256)	590080	activation_31[0][0]
bn4d_branch2b (BatchNormalizati	(None, 14, 14, 256)	1024	res4d_branch2b[0][0]

activation_32 (Activation)	(None, 14, 14, 256)	0	bn4d_branch2b[0][0]
res4d_branch2c (Conv2D)	(None, 14, 14, 1024)	263168	activation_32[0][0]
bn4d_branch2c (BatchNormalizati	(None, 14, 14, 1024)	4096	res4d_branch2c[0][0]
add_10 (Add)	(None, 14, 14, 1024)	0	bn4d_branch2c[0][0] activation_30[0][0]
activation_33 (Activation)	(None, 14, 14, 1024)	0	add_10[0][0]
res4e_branch2a (Conv2D)	(None, 14, 14, 256)	262400	activation_33[0][0]
bn4e_branch2a (BatchNormalizati	(None, 14, 14, 256)	1024	res4e_branch2a[0][0]
activation_34 (Activation)	(None, 14, 14, 256)	0	bn4e_branch2a[0][0]
res4e_branch2b (Conv2D)	(None, 14, 14, 256)	590080	activation_34[0][0]
bn4e_branch2b (BatchNormalizati	(None, 14, 14, 256)	1024	res4e_branch2b[0][0]
activation_35 (Activation)	(None, 14, 14, 256)	0	bn4e_branch2b[0][0]
res4e_branch2c (Conv2D)	(None, 14, 14, 1024)	263168	activation_35[0][0]
bn4e_branch2c (BatchNormalizati	(None, 14, 14, 1024)	4096	res4e_branch2c[0][0]
add_11 (Add)	(None, 14, 14, 1024)	0	bn4e_branch2c[0][0] activation_33[0][0]

activation_36 (Activation)	(None, 14, 14, 1024)	0	add_11[0][0]
res4f_branch2a (Conv2D)	(None, 14, 14, 256)	262400	activation_36[0][0]
bn4f_branch2a (BatchNormalizati	(None, 14, 14, 256)	1024	res4f_branch2a[0][0]
activation_37 (Activation)	(None, 14, 14, 256)	0	bn4f_branch2a[0][0]
res4f_branch2b (Conv2D)	(None, 14, 14, 256)	590080	activation_37[0][0]
bn4f_branch2b (BatchNormalizati	(None, 14, 14, 256)	1024	res4f_branch2b[0][0]
activation_38 (Activation)	(None, 14, 14, 256)	0	bn4f_branch2b[0][0]
res4f_branch2c (Conv2D)	(None, 14, 14, 1024)	263168	activation_38[0][0]
bn4f_branch2c (BatchNormalizati	(None, 14, 14, 1024)	4096	res4f_branch2c[0][0]
add_12 (Add)	(None, 14, 14, 1024)	0	bn4f_branch2c[0][0] activation_36[0][0]
activation_39 (Activation)	(None, 14, 14, 1024)	0	add_12[0][0]
res5a_branch2a (Conv2D)	(None, 7, 7, 512)	524800	activation_39[0][0]
bn5a_branch2a (BatchNormalizati	(None, 7, 7, 512)	2048	res5a_branch2a[0][0]
activation_40 (Activation)	(None, 7, 7, 512)	0	bn5a_branch2a[0][0]
res5a_branch2b (Conv2D)	(None, 7, 7, 512)	2359808	activation_40[0][0]

bn5a_branch2b (BatchNormalizati	(None, 7, 7, 512)	2048	res5a_branch2b[0][0]
activation_41 (Activation)	(None, 7, 7, 512)	0	bn5a_branch2b[0][0]
res5a_branch2c (Conv2D)	(None, 7, 7, 2048)	1050624	activation_41[0][0]
res5a_branch1 (Conv2D)	(None, 7, 7, 2048)	2099200	activation_39[0][0]
bn5a_branch2c (BatchNormalizati	(None, 7, 7, 2048)	8192	res5a_branch2c[0][0]
bn5a_branch1 (BatchNormalizatio	(None, 7, 7, 2048)	8192	res5a_branch1[0][0]
add_13 (Add)	(None, 7, 7, 2048)	0	bn5a_branch2c[0][0] bn5a_branch1[0][0]
activation_42 (Activation)	(None, 7, 7, 2048)	0	add_13[0][0]
res5b_branch2a (Conv2D)	(None, 7, 7, 512)	1049088	activation_42[0][0]
bn5b_branch2a (BatchNormalizati	(None, 7, 7, 512)	2048	res5b_branch2a[0][0]
activation_43 (Activation)	(None, 7, 7, 512)	0	bn5b_branch2a[0][0]
res5b_branch2b (Conv2D)	(None, 7, 7, 512)	2359808	activation_43[0][0]
bn5b_branch2b (BatchNormalizati	(None, 7, 7, 512)	2048	res5b_branch2b[0][0]
activation_44 (Activation)	(None, 7, 7, 512)	0	bn5b_branch2b[0][0]
res5b_branch2c (Conv2D)	(None, 7, 7, 2048)	1050624	activation_44[0][0]

bn5b_branch2c (BatchNormalizati	(None, 7, 7, 2048)	8192	res5b_branch2c[0][0]
add_14 (Add)	(None, 7, 7, 2048)	0	bn5b_branch2c[0][0] activation_42[0][0]
activation_45 (Activation)	(None, 7, 7, 2048)	0	add_14[0][0]
res5c_branch2a (Conv2D)	(None, 7, 7, 512)	1049088	activation_45[0][0]
bn5c_branch2a (BatchNormalizati	(None, 7, 7, 512)	2048	res5c_branch2a[0][0]
activation_46 (Activation)	(None, 7, 7, 512)	0	bn5c_branch2a[0][0]
res5c_branch2b (Conv2D)	(None, 7, 7, 512)	2359808	activation_46[0][0]
bn5c_branch2b (BatchNormalizati	(None, 7, 7, 512)	2048	res5c_branch2b[0][0]
activation_47 (Activation)	(None, 7, 7, 512)	0	bn5c_branch2b[0][0]
res5c_branch2c (Conv2D)	(None, 7, 7, 2048)	1050624	activation_47[0][0]
bn5c_branch2c (BatchNormalizati	(None, 7, 7, 2048)	8192	res5c_branch2c[0][0]
add_15 (Add)	(None, 7, 7, 2048)	0	bn5c_branch2c[0][0] activation_45[0][0]
activation_48 (Activation)	(None, 7, 7, 2048)	0	add_15[0][0]
average_pooling2d (AveragePooli	(None, 4, 4, 2048)	0	activation_48[0][0]
dropout (Dropout)	(None, 4, 4, 2048)	0	average_pooling2d[0][0]

```

-----
flatten (Flatten)                (None, 32768)      0      dropout[0][0]
-----
dense (Dense)                    (None, 512)        16777728  flatten[0][0]
-----
dense_1 (Dense)                  (None, 256)        131328    dense[0][0]
-----
dense_2 (Dense)                  (None, 3)          771       dense_1[0][0]
=====
Total params: 40,497,539
Trainable params: 40,444,419
Non-trainable params: 53,120
-----

```

```

2021-10-20 17:36:31.922874: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the ML
IR Optimization Passes are enabled (registered 2)

```

```

Epoch 1/100

```



```

2021-10-20 17:36:40.160462: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369] Loaded cuDNN version 8005

```

189/189 [=====] - 262s 1s/step - loss: 4.8214 - accuracy: 0.5751 - val_loss: 0.8612 - val_accuracy: 0.7094
Epoch 2/100
189/189 [=====] - 82s 430ms/step - loss: 0.9222 - accuracy: 0.6588 - val_loss: 1.4347 - val_accuracy: 0.2965
Epoch 3/100
189/189 [=====] - 82s 433ms/step - loss: 0.7233 - accuracy: 0.7029 - val_loss: 0.6663 - val_accuracy: 0.7094
Epoch 4/100
189/189 [=====] - 83s 440ms/step - loss: 0.5386 - accuracy: 0.7776 - val_loss: 0.7174 - val_accuracy: 0.6396
Epoch 5/100
189/189 [=====] - 85s 449ms/step - loss: 0.4483 - accuracy: 0.8256 - val_loss: 3.6309 - val_accuracy: 0.6908
Epoch 6/100
189/189 [=====] - 87s 457ms/step - loss: 0.4173 - accuracy: 0.8313 - val_loss: 0.5777 - val_accuracy: 0.7507
Epoch 7/100
189/189 [=====] - 83s 440ms/step - loss: 0.4813 - accuracy: 0.8074 - val_loss: 0.7154 - val_accuracy: 0.7174
Epoch 8/100
189/189 [=====] - 85s 449ms/step - loss: 0.2863 - accuracy: 0.8858 - val_loss: 0.7330 - val_accuracy: 0.7317
Epoch 9/100
189/189 [=====] - 93s 491ms/step - loss: 0.1337 - accuracy: 0.9506 - val_loss: 0.9259 - val_accuracy: 0.5904
Epoch 10/100
189/189 [=====] - 83s 438ms/step - loss: 0.1934 - accuracy: 0.9308 - val_loss: 1.2001 - val_accuracy: 0.6609
Epoch 11/100


```
189/189 [=====] - 81s 427ms/step - loss: 0.0668 - accuracy: 0.9806 - val_loss: 1.
0974 - val_accuracy: 0.7078
Epoch 12/100
189/189 [=====] - 83s 439ms/step - loss: 0.1214 - accuracy: 0.9676 - val_loss: 1.
0696 - val_accuracy: 0.6642
Epoch 13/100
189/189 [=====] - 82s 434ms/step - loss: 0.3060 - accuracy: 0.9096 - val_loss: 3.
2787 - val_accuracy: 0.5027
Epoch 14/100
189/189 [=====] - 83s 439ms/step - loss: 0.3663 - accuracy: 0.8594 - val_loss: 0.
9287 - val_accuracy: 0.6380
Epoch 15/100
189/189 [=====] - 84s 444ms/step - loss: 0.1385 - accuracy: 0.9529 - val_loss: 1.
9083 - val_accuracy: 0.7028
Epoch 16/100
189/189 [=====] - 85s 448ms/step - loss: 0.4247 - accuracy: 0.8367 - val_loss: 0.
6992 - val_accuracy: 0.6842
Epoch 00016: early stopping
```




Model saved successfully!

Confusion matrix for train data:

190/190 [=====] - 59s 312ms/step - loss: 0.6621 - accuracy: 0.6795

Score = [0.6620919108390808, 0.6794523000717163]

Accuracy = 0.6794522807885837



Confusion matrix for val/test data:

48/48 [=====] - 15s 305ms/step - loss: 0.7040 - accuracy: 0.6822

Score = [0.7039801478385925, 0.6821781992912292]

Accuracy = 0.6821782178217822

