

In [1]:

```
"""
Image Super-Resolution using Convolutional
Autoencoders
Author: Amruth Karun M V
Date: 06-Nov-2021
"""

import cv2
import numpy as np
import random
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from tensorflow.keras import Model, Input, regularizers
from tensorflow.keras.layers import (
    Dense, Conv2D, MaxPool2D,
    UpSampling2D, Add)
from tensorflow.keras.callbacks import EarlyStopping
from keras.preprocessing import image
import glob
from tqdm import tqdm
import matplotlib.pyplot as plt
import warnings;
warnings.filterwarnings('ignore')

INPUT_PATH = '../input/cifar10/cifar10_sample/'

def show_sample_image():
    """
    Displays original 32 x 32 images.
    Arguments: None
    Returns: Displays the image
    """

    cifar_sample = glob.glob(INPUT_PATH + '*.png')
    random_index = random.randint(0, 19)
    print("Image: ", random_index)
    img_path = cifar_sample[random_index]
    img = cv2.imread(img_path)
    plt.imshow(img)

def load_images():
    """
```

Loads sample images from cifar10 dataset.

2 images are taken from each class. Total 20 images.

Arguments: None

Returns: Train and val images

"""

```
cifar_sample = glob.glob(INPUT_PATH + '*.png')
print("Total images = ", len(cifar_sample))
all_images = []
for i in tqdm(cifar_sample):
    img = image.load_img(i, target_size=(32,32,3))
    img = image.img_to_array(img)
    img = img/255.
    all_images.append(img)
all_images = np.array(all_images)
train_x, val_x = train_test_split(all_images, random_state=32, test_size=0.2
)
return train_x, val_x
```

```
def pixalate_image(image, scale_percent = 50):
```

"""

*Lower the resolution of input image without
reducing the size*

Arguments:

image -- input image

scale_percent -- amount to be reduced

Returns: Pixalated image

"""

```
width = int(image.shape[1] * scale_percent / 100)
height = int(image.shape[0] * scale_percent / 100)
dim = (width, height)
small_image = cv2.resize(image, dim, interpolation = cv2.INTER_AREA)
width = int(small_image.shape[1] * 100 / scale_percent)
height = int(small_image.shape[0] * 100 / scale_percent)
dim = (width, height)
low_res_image = cv2.resize(small_image, dim, interpolation = cv2.INTER_AREA)
return low_res_image
```

```
def get_low_res_image(train_x, val_x):
```

"""

*Get low resolution images for train
and validation set*

Arguments:

```

    train_x  -- train set
    val_x    -- validation set

```

Returns: Low resolution data

```

"""

```

```

# get low resolution images for the train set

```

```

train_x_px = []

```

```

for i in range(train_x.shape[0]):
    temp = pixalate_image(train_x[i, :, :, :])
    train_x_px.append(temp)
train_x_px = np.array(train_x_px)

```

```

# get low resolution images for the validation set

```

```

val_x_px = []

```

```

for i in range(val_x.shape[0]):
    temp = pixalate_image(val_x[i, :, :, :])
    val_x_px.append(temp)
val_x_px = np.array(val_x_px)

```

```

return train_x_px, val_x_px

```

```

def train_model(train_x_px, val_x_px):

```

```

    """

```

Trains the Autoencoder network

Arguments:

```

    train_x_px  -- low resolution train set
    val_x_px    -- low resolution validation set

```

Returns: Autoencoder Network

```

    """

```

```

    Input_img = Input(shape=(32, 32, 3))

```

```

    #encoding architecture

```

```

    x1 = Conv2D(64, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l1(10e-10))(Input_img)

```

```

    x2 = Conv2D(64, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l1(10e-10))(x1)

```

```

    x3 = MaxPool2D(padding='same')(x2)

```

```

    x4 = Conv2D(128, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l1(10e-10))(x3)

```

```

    x5 = Conv2D(128, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l1(10e-10))(x4)

```

```

    x6 = MaxPool2D(padding='same')(x5)

```

```

    encoded = Conv2D(256, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l1(10e-10))(x6)

```

```

# decoding architecture
x7 = UpSampling2D()(encoded)
x8 = Conv2D(128, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l1(10e-10))(x7)
x9 = Conv2D(128, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l1(10e-10))(x8)
x10 = Add()([x5, x9])
x11 = UpSampling2D()(x10)
x12 = Conv2D(64, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l1(10e-10))(x11)
x13 = Conv2D(64, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l1(10e-10))(x12)
x14 = Add()([x2, x13])
decoded = Conv2D(3, (3, 3), padding='same', activation='relu', kernel_regularizer=regularizers.l1(10e-10))(x14)
autoencoder = Model(Input_img, decoded)
autoencoder.compile(optimizer='adam', loss='mse', metrics=['accuracy'])

autoencoder.summary()

early_stopper = EarlyStopping(monitor='val_loss', min_delta=0.0001, patience=8, verbose=1, mode='auto')
history = autoencoder.fit(train_x_px, train_x,
                          epochs=256,
                          batch_size=64,
                          shuffle=True,
                          validation_data=(val_x_px, val_x),
                          callbacks=[early_stopper])

autoencoder.save_weights("autoencoder.h5")
return autoencoder

def get_results(autoencoder, val_x, val_x_px):
    """
    Evaluate the autoencoder model using
    validation data and predict results
    Arguments:
        autoencoder    -- trained autoencoder model
        val_x          -- original validation data
        val_x_px       -- low resolution validation data
    """

    results = autoencoder.evaluate(val_x_px, val_x)
    print('val_loss = {}, val_accuracy = {}'.format(results[0], results[1]))

```

```
predictions = autoencoder.predict(val_x_px)

n = 4
plt.figure(figsize= (20,10))
for i in range(n):
    ax1 = plt.subplot(3, n, i+1)
    plt.imshow(val_x_px[i])
    ax1.get_xaxis().set_visible(False)
    ax1.get_yaxis().set_visible(False)
    ax1.title.set_text('Pixelated Image')

    ax2 = plt.subplot(3, n, i+1+n)
    plt.imshow(predictions[i])
    ax2.get_xaxis().set_visible(False)
    ax2.get_yaxis().set_visible(False)
    ax2.title.set_text('Predicted Image')

plt.show()
```

In [2]:

```
# Show 1 random image from train set
# show_sample_image()
train_x, val_x = load_images()
train_x_px, val_x_px = get_low_res_image(train_x, val_x)

# Train the model
autoencoder = train_model(train_x_px, val_x_px)
get_results(autoencoder, val_x, val_x_px)
```

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 32, 32, 3)]	0	
conv2d (Conv2D)	(None, 32, 32, 64)	1792	input_1[0][0]
conv2d_1 (Conv2D)	(None, 32, 32, 64)	36928	conv2d[0][0]
max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0	conv2d_1[0][0]
conv2d_2 (Conv2D)	(None, 16, 16, 128)	73856	max_pooling2d[0][0]
conv2d_3 (Conv2D)	(None, 16, 16, 128)	147584	conv2d_2[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 128)	0	conv2d_3[0][0]
conv2d_4 (Conv2D)	(None, 8, 8, 256)	295168	max_pooling2d_1[0][0]
up_sampling2d (UpSampling2D)	(None, 16, 16, 256)	0	conv2d_4[0][0]
conv2d_5 (Conv2D)	(None, 16, 16, 128)	295040	up_sampling2d[0][0]

conv2d_6 (Conv2D)	(None, 16, 16, 128)	147584	conv2d_5[0][0]
add (Add)	(None, 16, 16, 128)	0	conv2d_3[0][0]
			conv2d_6[0][0]
up_sampling2d_1 (UpSampling2D)	(None, 32, 32, 128)	0	add[0][0]
conv2d_7 (Conv2D)	(None, 32, 32, 64)	73792	up_sampling2d_1[0][0]
conv2d_8 (Conv2D)	(None, 32, 32, 64)	36928	conv2d_7[0][0]
add_1 (Add)	(None, 32, 32, 64)	0	conv2d_1[0][0]
			conv2d_8[0][0]
conv2d_9 (Conv2D)	(None, 32, 32, 3)	1731	add_1[0][0]
Total params: 1,110,403			
Trainable params: 1,110,403			
Non-trainable params: 0			
Epoch 1/256			




```
2021-11-06 14:56:09.485213: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the MLIR Optimization Passes are enabled (registered 2)
```

1/1 [=====] - 2s 2s/step - loss: 0.2602 -
accuracy: 0.1879 - val_loss: 0.2291 - val_accuracy: 0.3660
Epoch 2/256
1/1 [=====] - 0s 335ms/step - loss: 0.2234
- accuracy: 0.1685 - val_loss: 0.1495 - val_accuracy: 0.3660
Epoch 3/256
1/1 [=====] - 0s 334ms/step - loss: 0.1438
- accuracy: 0.1685 - val_loss: 0.0832 - val_accuracy: 0.3667
Epoch 4/256
1/1 [=====] - 0s 335ms/step - loss: 0.0877
- accuracy: 0.1703 - val_loss: 0.0470 - val_accuracy: 0.3765
Epoch 5/256
1/1 [=====] - 0s 348ms/step - loss: 0.0555
- accuracy: 0.2001 - val_loss: 0.0368 - val_accuracy: 0.5352
Epoch 6/256
1/1 [=====] - 0s 337ms/step - loss: 0.0418
- accuracy: 0.5106 - val_loss: 0.0262 - val_accuracy: 0.4985
Epoch 7/256
1/1 [=====] - 0s 336ms/step - loss: 0.0290
- accuracy: 0.5726 - val_loss: 0.0232 - val_accuracy: 0.4189
Epoch 8/256
1/1 [=====] - 0s 349ms/step - loss: 0.0240
- accuracy: 0.4974 - val_loss: 0.0246 - val_accuracy: 0.2380
Epoch 9/256
1/1 [=====] - 0s 335ms/step - loss: 0.0232
- accuracy: 0.3387 - val_loss: 0.0268 - val_accuracy: 0.1758
Epoch 10/256
1/1 [=====] - 0s 331ms/step - loss: 0.0236
- accuracy: 0.2985 - val_loss: 0.0283 - val_accuracy: 0.1641
Epoch 11/256
1/1 [=====] - 0s 336ms/step - loss: 0.0240
- accuracy: 0.2819 - val_loss: 0.0283 - val_accuracy: 0.1633
Epoch 12/256
1/1 [=====] - 0s 336ms/step - loss: 0.0236
- accuracy: 0.2809 - val_loss: 0.0271 - val_accuracy: 0.1641
Epoch 13/256
1/1 [=====] - 0s 330ms/step - loss: 0.0227
- accuracy: 0.2849 - val_loss: 0.0252 - val_accuracy: 0.1689
Epoch 14/256
1/1 [=====] - 0s 390ms/step - loss: 0.0215
- accuracy: 0.3007 - val_loss: 0.0230 - val_accuracy: 0.1807
Epoch 15/256
1/1 [=====] - 0s 337ms/step - loss: 0.0202
- accuracy: 0.3250 - val_loss: 0.0209 - val_accuracy: 0.2034

Epoch 16/256

1/1 [=====] - 0s 334ms/step - loss: 0.0191
- accuracy: 0.3649 - val_loss: 0.0191 - val_accuracy: 0.2349

Epoch 17/256

1/1 [=====] - 0s 335ms/step - loss: 0.0181
- accuracy: 0.4069 - val_loss: 0.0177 - val_accuracy: 0.2773

Epoch 18/256

1/1 [=====] - 0s 334ms/step - loss: 0.0174
- accuracy: 0.4540 - val_loss: 0.0165 - val_accuracy: 0.3789

Epoch 19/256

1/1 [=====] - 0s 333ms/step - loss: 0.0166
- accuracy: 0.4990 - val_loss: 0.0156 - val_accuracy: 0.4431

Epoch 20/256

1/1 [=====] - 0s 334ms/step - loss: 0.0159
- accuracy: 0.5351 - val_loss: 0.0145 - val_accuracy: 0.4714

Epoch 21/256

1/1 [=====] - 0s 330ms/step - loss: 0.0152
- accuracy: 0.5490 - val_loss: 0.0135 - val_accuracy: 0.4795

Epoch 22/256

1/1 [=====] - 0s 334ms/step - loss: 0.0144
- accuracy: 0.5515 - val_loss: 0.0125 - val_accuracy: 0.4785

Epoch 23/256

1/1 [=====] - 0s 332ms/step - loss: 0.0135
- accuracy: 0.5461 - val_loss: 0.0117 - val_accuracy: 0.4521

Epoch 24/256

1/1 [=====] - 0s 335ms/step - loss: 0.0126
- accuracy: 0.5223 - val_loss: 0.0111 - val_accuracy: 0.3860

Epoch 25/256

1/1 [=====] - 0s 334ms/step - loss: 0.0119
- accuracy: 0.4785 - val_loss: 0.0102 - val_accuracy: 0.3877

Epoch 26/256

1/1 [=====] - 0s 336ms/step - loss: 0.0111
- accuracy: 0.4750 - val_loss: 0.0093 - val_accuracy: 0.4441

Epoch 27/256

1/1 [=====] - 0s 334ms/step - loss: 0.0102
- accuracy: 0.4880 - val_loss: 0.0085 - val_accuracy: 0.4929

Epoch 28/256

1/1 [=====] - 0s 335ms/step - loss: 0.0094
- accuracy: 0.5021 - val_loss: 0.0074 - val_accuracy: 0.5217

Epoch 29/256

1/1 [=====] - 0s 333ms/step - loss: 0.0086
- accuracy: 0.5152 - val_loss: 0.0065 - val_accuracy: 0.5698

Epoch 30/256

1/1 [=====] - 0s 334ms/step - loss: 0.0079
- accuracy: 0.5339 - val_loss: 0.0059 - val_accuracy: 0.6138

Epoch 31/256

1/1 [=====] - 0s 333ms/step - loss: 0.0073
- accuracy: 0.5491 - val_loss: 0.0056 - val_accuracy: 0.6296

Epoch 32/256

1/1 [=====] - 0s 333ms/step - loss: 0.0068
- accuracy: 0.5527 - val_loss: 0.0059 - val_accuracy: 0.5566

Epoch 33/256

1/1 [=====] - 0s 333ms/step - loss: 0.0068
- accuracy: 0.5243 - val_loss: 0.0059 - val_accuracy: 0.6985

Epoch 34/256

1/1 [=====] - 0s 341ms/step - loss: 0.0073
- accuracy: 0.5950 - val_loss: 0.0050 - val_accuracy: 0.6978

Epoch 35/256

1/1 [=====] - 0s 338ms/step - loss: 0.0060
- accuracy: 0.5881 - val_loss: 0.0057 - val_accuracy: 0.6379

Epoch 36/256

1/1 [=====] - 0s 331ms/step - loss: 0.0064
- accuracy: 0.5712 - val_loss: 0.0051 - val_accuracy: 0.7119

Epoch 37/256

1/1 [=====] - 0s 333ms/step - loss: 0.0062
- accuracy: 0.6049 - val_loss: 0.0048 - val_accuracy: 0.6780

Epoch 38/256

1/1 [=====] - 0s 339ms/step - loss: 0.0058
- accuracy: 0.5925 - val_loss: 0.0052 - val_accuracy: 0.5488

Epoch 39/256

1/1 [=====] - 0s 335ms/step - loss: 0.0058
- accuracy: 0.5678 - val_loss: 0.0047 - val_accuracy: 0.6365

Epoch 40/256

1/1 [=====] - 0s 333ms/step - loss: 0.0054
- accuracy: 0.5986 - val_loss: 0.0047 - val_accuracy: 0.7190

Epoch 41/256

1/1 [=====] - 0s 335ms/step - loss: 0.0054
- accuracy: 0.6285 - val_loss: 0.0047 - val_accuracy: 0.7109

Epoch 42/256

1/1 [=====] - 0s 335ms/step - loss: 0.0053
- accuracy: 0.6296 - val_loss: 0.0045 - val_accuracy: 0.6477

Epoch 43/256

1/1 [=====] - 0s 336ms/step - loss: 0.0051
- accuracy: 0.6281 - val_loss: 0.0047 - val_accuracy: 0.5620

Epoch 44/256

1/1 [=====] - 0s 337ms/step - loss: 0.0052
- accuracy: 0.6267 - val_loss: 0.0045 - val_accuracy: 0.5894

Epoch 45/256

1/1 [=====] - 0s 338ms/step - loss: 0.0049
- accuracy: 0.6432 - val_loss: 0.0046 - val_accuracy: 0.6704

Epoch 46/256

1/1 [=====] - 0s 381ms/step - loss: 0.0050
- accuracy: 0.6653 - val_loss: 0.0043 - val_accuracy: 0.7063

Epoch 47/256

1/1 [=====] - 0s 338ms/step - loss: 0.0048
- accuracy: 0.6780 - val_loss: 0.0043 - val_accuracy: 0.6970

Epoch 48/256

1/1 [=====] - 0s 336ms/step - loss: 0.0047
- accuracy: 0.6880 - val_loss: 0.0043 - val_accuracy: 0.6350

Epoch 49/256

1/1 [=====] - 0s 334ms/step - loss: 0.0046
- accuracy: 0.6982 - val_loss: 0.0043 - val_accuracy: 0.5803

Epoch 50/256

1/1 [=====] - 0s 337ms/step - loss: 0.0046
- accuracy: 0.6943 - val_loss: 0.0042 - val_accuracy: 0.6528

Epoch 51/256

1/1 [=====] - 0s 334ms/step - loss: 0.0045
- accuracy: 0.7189 - val_loss: 0.0041 - val_accuracy: 0.7188

Epoch 52/256

1/1 [=====] - 0s 333ms/step - loss: 0.0044
- accuracy: 0.7158 - val_loss: 0.0040 - val_accuracy: 0.7253

Epoch 53/256

1/1 [=====] - 0s 335ms/step - loss: 0.0043
- accuracy: 0.7310 - val_loss: 0.0040 - val_accuracy: 0.6611

Epoch 54/256

1/1 [=====] - 0s 331ms/step - loss: 0.0042
- accuracy: 0.7432 - val_loss: 0.0041 - val_accuracy: 0.6228

Epoch 55/256

1/1 [=====] - 0s 335ms/step - loss: 0.0042
- accuracy: 0.7355 - val_loss: 0.0039 - val_accuracy: 0.6990

Epoch 56/256

1/1 [=====] - 0s 335ms/step - loss: 0.0041
- accuracy: 0.7606 - val_loss: 0.0039 - val_accuracy: 0.7405

Epoch 57/256

1/1 [=====] - 0s 334ms/step - loss: 0.0041
- accuracy: 0.7505 - val_loss: 0.0038 - val_accuracy: 0.7305

Epoch 58/256

1/1 [=====] - 0s 334ms/step - loss: 0.0040
- accuracy: 0.7729 - val_loss: 0.0039 - val_accuracy: 0.6655

Epoch 59/256

1/1 [=====] - 0s 333ms/step - loss: 0.0040
- accuracy: 0.7598 - val_loss: 0.0038 - val_accuracy: 0.7219

Epoch 60/256

1/1 [=====] - 0s 331ms/step - loss: 0.0038
- accuracy: 0.7805 - val_loss: 0.0038 - val_accuracy: 0.7646

Epoch 61/256

1/1 [=====] - 0s 337ms/step - loss: 0.0038
- accuracy: 0.7836 - val_loss: 0.0038 - val_accuracy: 0.7698

Epoch 62/256

1/1 [=====] - 0s 331ms/step - loss: 0.0038
- accuracy: 0.7958 - val_loss: 0.0038 - val_accuracy: 0.7314

Epoch 63/256

1/1 [=====] - 0s 339ms/step - loss: 0.0038
- accuracy: 0.7844 - val_loss: 0.0038 - val_accuracy: 0.8025

Epoch 64/256

1/1 [=====] - 0s 346ms/step - loss: 0.0037
- accuracy: 0.8094 - val_loss: 0.0037 - val_accuracy: 0.8040

Epoch 65/256

1/1 [=====] - 0s 333ms/step - loss: 0.0037
- accuracy: 0.8049 - val_loss: 0.0037 - val_accuracy: 0.7686

Epoch 66/256

1/1 [=====] - 0s 332ms/step - loss: 0.0037
- accuracy: 0.7899 - val_loss: 0.0037 - val_accuracy: 0.7817

Epoch 67/256

1/1 [=====] - 0s 335ms/step - loss: 0.0036
- accuracy: 0.7948 - val_loss: 0.0037 - val_accuracy: 0.8113

Epoch 68/256

1/1 [=====] - 0s 337ms/step - loss: 0.0036
- accuracy: 0.8069 - val_loss: 0.0037 - val_accuracy: 0.8123

Epoch 69/256

1/1 [=====] - 0s 331ms/step - loss: 0.0036
- accuracy: 0.8096 - val_loss: 0.0036 - val_accuracy: 0.7820

Epoch 70/256

1/1 [=====] - 0s 334ms/step - loss: 0.0036
- accuracy: 0.8033 - val_loss: 0.0036 - val_accuracy: 0.7717

Epoch 71/256

1/1 [=====] - 0s 333ms/step - loss: 0.0035
- accuracy: 0.8011 - val_loss: 0.0036 - val_accuracy: 0.7888

Epoch 72/256

1/1 [=====] - 0s 334ms/step - loss: 0.0035
- accuracy: 0.8110 - val_loss: 0.0036 - val_accuracy: 0.7954

Epoch 73/256

1/1 [=====] - 0s 339ms/step - loss: 0.0035
- accuracy: 0.8140 - val_loss: 0.0036 - val_accuracy: 0.7791

Epoch 74/256

1/1 [=====] - 0s 335ms/step - loss: 0.0035
- accuracy: 0.8094 - val_loss: 0.0036 - val_accuracy: 0.7554

Epoch 75/256

1/1 [=====] - 0s 332ms/step - loss: 0.0035
- accuracy: 0.8044 - val_loss: 0.0035 - val_accuracy: 0.7815

Epoch 76/256

1/1 [=====] - 0s 333ms/step - loss: 0.0034
- accuracy: 0.8096 - val_loss: 0.0035 - val_accuracy: 0.7986

Epoch 77/256

1/1 [=====] - 0s 333ms/step - loss: 0.0034
- accuracy: 0.8141 - val_loss: 0.0035 - val_accuracy: 0.7837

Epoch 78/256

1/1 [=====] - 0s 333ms/step - loss: 0.0034
- accuracy: 0.8115 - val_loss: 0.0035 - val_accuracy: 0.7712

Epoch 79/256

1/1 [=====] - 0s 394ms/step - loss: 0.0034
- accuracy: 0.8088 - val_loss: 0.0035 - val_accuracy: 0.7927

Epoch 80/256

1/1 [=====] - 0s 335ms/step - loss: 0.0033
- accuracy: 0.8152 - val_loss: 0.0035 - val_accuracy: 0.8062

Epoch 81/256

1/1 [=====] - 0s 330ms/step - loss: 0.0033
- accuracy: 0.8239 - val_loss: 0.0035 - val_accuracy: 0.7981

Epoch 82/256

1/1 [=====] - 0s 364ms/step - loss: 0.0033
- accuracy: 0.8149 - val_loss: 0.0035 - val_accuracy: 0.7903

Epoch 83/256

1/1 [=====] - 0s 335ms/step - loss: 0.0033
- accuracy: 0.8110 - val_loss: 0.0035 - val_accuracy: 0.8022

Epoch 84/256

1/1 [=====] - 0s 332ms/step - loss: 0.0033
- accuracy: 0.8178 - val_loss: 0.0035 - val_accuracy: 0.8096

Epoch 85/256

1/1 [=====] - 0s 335ms/step - loss: 0.0033
- accuracy: 0.8262 - val_loss: 0.0034 - val_accuracy: 0.8110

Epoch 86/256

1/1 [=====] - 0s 333ms/step - loss: 0.0032
- accuracy: 0.8148 - val_loss: 0.0034 - val_accuracy: 0.8013

Epoch 87/256

1/1 [=====] - 0s 332ms/step - loss: 0.0032
- accuracy: 0.8125 - val_loss: 0.0034 - val_accuracy: 0.8159

Epoch 88/256

1/1 [=====] - 0s 338ms/step - loss: 0.0032
- accuracy: 0.8238 - val_loss: 0.0034 - val_accuracy: 0.8159

Epoch 00088: early stopping

1/1 [=====] - 0s 59ms/step - loss: 0.0034
- accuracy: 0.8159

val_loss = 0.0034120238851755857, val_accuracy = 0.81591796875

Pixelated Image



Pixelated Image




Pixelated Image




Pixelated Image




Predicted Image



Predicted Image



Predicted Image



Predicted Image

