

In [1]:

```
"""
Tumor Classssification using VGG-19
Author: Amruth Karun M V
Date: 07-Nov-2021
"""

import os
import cv2
import numpy as np
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.models import Model
from tensorflow.keras.layers import (
    Dense, Input, Dropout, Flatten,
    Conv2D, MaxPooling2D, BatchNormalization
)

from sklearn import metrics
import matplotlib.pyplot as plt
%matplotlib inline

TRAIN_PATH = "../input/brain-tumor-mri-dataset/Training"
TEST_PATH = "../input/brain-tumor-mri-dataset/Testing"
CLASS_NAMES = ['Glioma', 'Meningioma', 'No-tumor', 'Pituitary']
EPOCHS = 100
BATCH_SIZE = 128
LEARNING_RATE = 0.001
```

```

def plot_sample_images():
    """
    Plots sample images for each class
    Arguments: None
    Returns: Plots sample data
    """

    plt.figure(figsize=(10, 10))
    sample_image_path = ['/glioma/Tr-gl_0010.jpg', '/meningioma/Tr-me_0010.jpg',
                        '/notumor/Tr-no_0010.jpg', '/pituitary/Tr-pi_0010.jpg']
    for i in range(len(CLASS_NAMES)):
        ax = plt.subplot(2, 2, i + 1)
        img = cv2.imread(TRAIN_PATH + sample_image_path[i])
        img = cv2.resize(img, (128, 128))
        plt.imshow(img)
        plt.title(CLASS_NAMES[i])

def load_data(input_path, shuffle=False):
    """
    Loads input data fro directory
    Arguments:
        input_path -- input data path
        shuffle    -- whether data needs to be shuffled or not
    Returns: Data generator
    """

    data_generator = keras.preprocessing.image.ImageDataGenerator()
    data_generator = data_generator.flow_from_directory(directory=input_path, target_size=(224,224),
                                                    shuffle=shuffle, class_mode="categorical")

```

```
return data_generator
```

```
def load_model():
```

```
    """
```

```
    Creates a keras VGG-19 model
```

```
    Arguments: None
```

```
    Returns: VGG-19 Model
```

```
    """
```

```
img_input = Input(shape=(224, 224, 3))
```

```
# Block 1
```

```
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv1')(img_input)
```

```
x = Conv2D(64, (3, 3), activation='relu', padding='same', name='block1_conv2')(x)
```

```
x = BatchNormalization()(x)
```

```
x = MaxPooling2D((2, 2), strides=(2, 2), name='block1_pool')(x)
```

```
# Block 2
```

```
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv1')(x)
```

```
x = Conv2D(128, (3, 3), activation='relu', padding='same', name='block2_conv2')(x)
```

```
x = BatchNormalization()(x)
```

```
x = MaxPooling2D((2, 2), strides=(2, 2), name='block2_pool')(x)
```

```
# Block 3
```

```
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv1')(x)
```

```
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv2')(x)
```

```
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv3')(x)
```

```
x = Conv2D(256, (3, 3), activation='relu', padding='same', name='block3_conv4')(x)
```

```
x = BatchNormalization()(x)
```

```
x = MaxPooling2D((2, 2), strides=(2, 2), name='block3_pool')(x)
```

```
# Block 4
```

```
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv3')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block4_conv4')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block4_pool')(x)
```

```
# Block 5
```

```
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv1')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv2')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv3')(x)
x = Conv2D(512, (3, 3), activation='relu', padding='same', name='block5_conv4')(x)
x = BatchNormalization()(x)
x = MaxPooling2D((2, 2), strides=(2, 2), name='block5_pool')(x)
```

```
# Classification block
```

```
x = Flatten(name='flatten')(x)
x = Dense(4096, activation='relu', name='fc1')(x)
x = Dropout(0.4)(x)
x = Dense(4096, activation='relu', name='fc2')(x)
x = Dropout(0.4)(x)
x = Dense(4, activation='softmax', name='predictions')(x)
```

```
model = Model(img_input, x, name='vgg-19')
```

```
model.summary()
```

```
opt = Adam(learning_rate=LEARNING_RATE)
```

```
model.compile(loss = keras.losses.categorical_crossentropy, optimizer=opt, metrics=['accuracy'])
```

```
return model
```

```
def plot_curves(history):
```

```
    """
```

```
    Plots loss and accuracy and loss plots for  
    training and validation datasets
```

```
    Arguments:
```

```
        history -- training history
```

```
    Returns: None
```

```
    """
```

```
    plt.plot(history.history['loss'], label="Training loss")
```

```
    plt.plot(history.history['val_loss'], label="Validation loss")
```

```
    plt.legend()
```

```
    plt.title('Training Loss VS Validation Loss')
```

```
    plt.show()
```

```
    plt.plot(history.history['accuracy'], label="Training accuracy")
```

```
    plt.plot(history.history['val_accuracy'], label="Validation accuracy")
```

```
    plt.title('Training Accuracy VS Validation Accuracy')
```

```
    plt.legend()
```

```
    plt.show()
```

```
def evaluate_model(model, input_path):
```

```
    """
```

```
    Evaluates the model and displays  
    the confusion matrix
```

```
    Arguments:
```

```

        model          -- trained model
        input_path   -- input data path
Returns: Model score and confusion matrix
"""

data_generator = load_data(input_path)
predictions = model.predict(data_generator, BATCH_SIZE)
y_pred = np.argmax(predictions, axis=1)
y_true = data_generator.classes

print("Score = ", model.evaluate(data_generator))
print("Accuracy = ", metrics.accuracy_score(y_true, y_pred))
cm = metrics.confusion_matrix(y_true, y_pred)
metrics.ConfusionMatrixDisplay(cm, display_labels=CLASS_NAMES).plot(cmap=plt.cm.Blues,
                                                                    xticks_rotation='vertical')

plt.show()

def train_model():
    """

    Trains VGG-19 model and saves the
    trained weights to an H5 file.
    Arguments: None
    Returns: None
    """

    train_generator = load_data(TRAIN_PATH, True)
    val_generator = load_data(TEST_PATH, True)

    # Loads VGG-19 model
    model = load_model()

```

```
earlystop = keras.callbacks.EarlyStopping(monitor='loss', min_delta=1e-11, patience=10)
reduce_lr = keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor=0.2,
                                              patience=6, verbose=1)

model_callbacks = [earlystop, reduce_lr]


history = model.fit(
    train_generator,
    batch_size=BATCH_SIZE,
    epochs=EPOCHS,
    validation_data=val_generator,
    validation_steps=val_generator.samples//BATCH_SIZE,
    steps_per_epoch=train_generator.samples//BATCH_SIZE,
    callbacks=model_callbacks)

plot_curves(history)
model.save_weights("model_vgg19.h5")
print("Model saved successfully!")

return model
```

In [2]:

```
plot_sample_images()
```





In [3]:

```
# Train the model
model = train_model()
print("Confusion matrix for train data: ")
evaluate_model(model, TRAIN_PATH)
print("Confusion matrix for val data: ")
evaluate_model(model, TEST_PATH)
```

Found 5712 images belonging to 4 classes.

Found 1311 images belonging to 4 classes.

Model: "vgg-19"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
batch_normalization (BatchNo	(None, 224, 224, 64)	256
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
batch_normalization_1 (Batch	(None, 112, 112, 128)	512
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv4 (Conv2D)	(None, 56, 56, 256)	590080

batch_normalization_2 (Batch Normalization)	(None, 56, 56, 256)	1024
-----		
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
-----		
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
-----		
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
-----		
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
-----		
block4_conv4 (Conv2D)	(None, 28, 28, 512)	2359808
-----		
batch_normalization_3 (Batch Normalization)	(None, 28, 28, 512)	2048
-----		
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
-----		
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
-----		
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
-----		
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
-----		
block5_conv4 (Conv2D)	(None, 14, 14, 512)	2359808
-----		
batch_normalization_4 (Batch Normalization)	(None, 14, 14, 512)	2048
-----		
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
-----		
flatten (Flatten)	(None, 25088)	0
-----		
fc1 (Dense)	(None, 4096)	102764544

```

-----
dropout (Dropout)          (None, 4096)          0
-----
fc2 (Dense)                (None, 4096)        16781312
-----
dropout_1 (Dropout)        (None, 4096)          0
-----
predictions (Dense)        (None, 4)            16388
=====
Total params: 139,592,516
Trainable params: 139,589,572
Non-trainable params: 2,944
-----

```

```

2021-11-07 08:13:08.730421: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:185] None of the ML
IR Optimization Passes are enabled (registered 2)

```

```

Epoch 1/100

```

```

2021-11-07 08:13:11.151784: I tensorflow/stream_executor/cuda/cuda_dnn.cc:369] Loaded cuDNN version 8005

```

44/44 [=====] - 27s 347ms/step - loss: 18.9244 - accuracy: 0.3956 - val\_loss: 107  
28.3447 - val\_accuracy: 0.3406  
Epoch 2/100  
44/44 [=====] - 14s 323ms/step - loss: 1.8303 - accuracy: 0.4780 - val\_loss: 67.6  
216 - val\_accuracy: 0.3156  
Epoch 3/100  
44/44 [=====] - 14s 322ms/step - loss: 1.6955 - accuracy: 0.4993 - val\_loss: 67.5  
509 - val\_accuracy: 0.3344  
Epoch 4/100  
44/44 [=====] - 16s 356ms/step - loss: 1.0784 - accuracy: 0.5402 - val\_loss: 4.29  
85 - val\_accuracy: 0.4313  
Epoch 5/100  
44/44 [=====] - 14s 309ms/step - loss: 0.9957 - accuracy: 0.5888 - val\_loss: 11.3  
010 - val\_accuracy: 0.3781  
Epoch 6/100  
44/44 [=====] - 14s 305ms/step - loss: 1.8594 - accuracy: 0.5192 - val\_loss: 2.80  
99 - val\_accuracy: 0.4594  
Epoch 7/100  
44/44 [=====] - 13s 304ms/step - loss: 1.1025 - accuracy: 0.5866 - val\_loss: 1.05  
70 - val\_accuracy: 0.5969  
Epoch 8/100  
44/44 [=====] - 13s 301ms/step - loss: 0.9950 - accuracy: 0.5724 - val\_loss: 0.87  
98 - val\_accuracy: 0.6875  
Epoch 9/100  
44/44 [=====] - 14s 304ms/step - loss: 1.2834 - accuracy: 0.5866 - val\_loss: 2.83  
24 - val\_accuracy: 0.4125  
Epoch 10/100  
44/44 [=====] - 13s 298ms/step - loss: 1.2852 - accuracy: 0.6406 - val\_loss: 1.12  
71 - val\_accuracy: 0.5219  
Epoch 11/100

44/44 [=====] - 13s 299ms/step - loss: 1.0293 - accuracy: 0.6243 - val\_loss: 1.27  
10 - val\_accuracy: 0.4906  
Epoch 12/100  
44/44 [=====] - 13s 298ms/step - loss: 0.9163 - accuracy: 0.6406 - val\_loss: 1.43  
12 - val\_accuracy: 0.5875  
Epoch 13/100  
44/44 [=====] - 13s 302ms/step - loss: 1.2838 - accuracy: 0.6151 - val\_loss: 1.17  
62 - val\_accuracy: 0.6531  
Epoch 14/100  
44/44 [=====] - 13s 298ms/step - loss: 0.8429 - accuracy: 0.6712 - val\_loss: 1.49  
01 - val\_accuracy: 0.4656

Epoch 00014: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 15/100  
44/44 [=====] - 13s 298ms/step - loss: 0.7949 - accuracy: 0.6746 - val\_loss: 0.81  
77 - val\_accuracy: 0.6906  
Epoch 16/100  
44/44 [=====] - 13s 303ms/step - loss: 0.7275 - accuracy: 0.7230 - val\_loss: 0.73  
94 - val\_accuracy: 0.6781  
Epoch 17/100  
44/44 [=====] - 13s 300ms/step - loss: 0.6466 - accuracy: 0.7141 - val\_loss: 0.65  
46 - val\_accuracy: 0.7250  
Epoch 18/100  
44/44 [=====] - 13s 300ms/step - loss: 0.6264 - accuracy: 0.7322 - val\_loss: 0.57  
77 - val\_accuracy: 0.7594  
Epoch 19/100  
44/44 [=====] - 13s 299ms/step - loss: 0.5872 - accuracy: 0.7450 - val\_loss: 0.66  
47 - val\_accuracy: 0.6938  
Epoch 20/100  
44/44 [=====] - 13s 294ms/step - loss: 0.6402 - accuracy: 0.7342 - val\_loss: 0.62  
28 - val\_accuracy: 0.7375

Epoch 21/100

44/44 [=====] - 13s 299ms/step - loss: 0.6015 - accuracy: 0.7607 - val\_loss: 0.6161 - val\_accuracy: 0.7219

Epoch 22/100

44/44 [=====] - 13s 303ms/step - loss: 0.5820 - accuracy: 0.7521 - val\_loss: 0.6057 - val\_accuracy: 0.7437

Epoch 23/100

44/44 [=====] - 13s 294ms/step - loss: 0.5488 - accuracy: 0.7708 - val\_loss: 0.5667 - val\_accuracy: 0.7250

Epoch 24/100

44/44 [=====] - 13s 297ms/step - loss: 0.5017 - accuracy: 0.8017 - val\_loss: 0.6481 - val\_accuracy: 0.7281

Epoch 25/100

44/44 [=====] - 13s 297ms/step - loss: 0.5340 - accuracy: 0.7692 - val\_loss: 0.5278 - val\_accuracy: 0.7937

Epoch 26/100

44/44 [=====] - 13s 299ms/step - loss: 0.5419 - accuracy: 0.7756 - val\_loss: 0.5529 - val\_accuracy: 0.7594

Epoch 27/100

44/44 [=====] - 14s 308ms/step - loss: 0.5835 - accuracy: 0.7692 - val\_loss: 0.7031 - val\_accuracy: 0.7000

Epoch 28/100

44/44 [=====] - 13s 298ms/step - loss: 0.5325 - accuracy: 0.7884 - val\_loss: 0.5754 - val\_accuracy: 0.7781

Epoch 29/100

44/44 [=====] - 13s 297ms/step - loss: 0.4807 - accuracy: 0.8118 - val\_loss: 0.5418 - val\_accuracy: 0.7688

Epoch 30/100

44/44 [=====] - 13s 300ms/step - loss: 0.4834 - accuracy: 0.7969 - val\_loss: 0.5349 - val\_accuracy: 0.7531

Epoch 31/100



44/44 [=====] - 13s 298ms/step - loss: 0.4834 - accuracy: 0.8011 - val\_loss: 0.52  
32 - val\_accuracy: 0.7781  
Epoch 32/100  
44/44 [=====] - 13s 301ms/step - loss: 0.4476 - accuracy: 0.8104 - val\_loss: 0.51  
57 - val\_accuracy: 0.7688  
Epoch 33/100  
44/44 [=====] - 13s 293ms/step - loss: 0.4554 - accuracy: 0.8089 - val\_loss: 0.57  
41 - val\_accuracy: 0.7656  
Epoch 34/100  
44/44 [=====] - 13s 296ms/step - loss: 0.4515 - accuracy: 0.8053 - val\_loss: 0.50  
34 - val\_accuracy: 0.7688  
Epoch 35/100  
44/44 [=====] - 13s 300ms/step - loss: 0.4619 - accuracy: 0.7976 - val\_loss: 0.54  
47 - val\_accuracy: 0.7719  
Epoch 36/100  
44/44 [=====] - 13s 297ms/step - loss: 0.4471 - accuracy: 0.8203 - val\_loss: 0.55  
81 - val\_accuracy: 0.7625  
Epoch 37/100  
44/44 [=====] - 13s 296ms/step - loss: 0.4316 - accuracy: 0.8267 - val\_loss: 0.48  
50 - val\_accuracy: 0.7812  
Epoch 38/100  
44/44 [=====] - 13s 297ms/step - loss: 0.4427 - accuracy: 0.8203 - val\_loss: 0.51  
10 - val\_accuracy: 0.7969  
Epoch 39/100  
44/44 [=====] - 13s 298ms/step - loss: 0.4593 - accuracy: 0.8189 - val\_loss: 0.51  
88 - val\_accuracy: 0.7750  
Epoch 40/100  
44/44 [=====] - 13s 299ms/step - loss: 0.4260 - accuracy: 0.8295 - val\_loss: 0.52  
90 - val\_accuracy: 0.7750  
Epoch 41/100  
44/44 [=====] - 13s 300ms/step - loss: 0.4570 - accuracy: 0.8210 - val\_loss: 0.45

01 - val\_accuracy: 0.8250  
Epoch 42/100  
44/44 [=====] - 13s 296ms/step - loss: 0.4136 - accuracy: 0.8345 - val\_loss: 0.44  
60 - val\_accuracy: 0.8125  
Epoch 43/100  
44/44 [=====] - 13s 296ms/step - loss: 0.4211 - accuracy: 0.8317 - val\_loss: 0.47  
14 - val\_accuracy: 0.8062  
Epoch 44/100  
44/44 [=====] - 13s 297ms/step - loss: 0.3930 - accuracy: 0.8317 - val\_loss: 0.44  
40 - val\_accuracy: 0.8125  
Epoch 45/100  
44/44 [=====] - 13s 297ms/step - loss: 0.4150 - accuracy: 0.8388 - val\_loss: 0.50  
96 - val\_accuracy: 0.7812  
Epoch 46/100  
44/44 [=====] - 13s 297ms/step - loss: 0.4074 - accuracy: 0.8303 - val\_loss: 0.44  
36 - val\_accuracy: 0.7906  
Epoch 47/100  
44/44 [=====] - 13s 298ms/step - loss: 0.3636 - accuracy: 0.8473 - val\_loss: 0.55  
71 - val\_accuracy: 0.7500  
Epoch 48/100  
44/44 [=====] - 13s 295ms/step - loss: 0.3985 - accuracy: 0.8551 - val\_loss: 0.47  
46 - val\_accuracy: 0.8031  
Epoch 49/100  
44/44 [=====] - 13s 302ms/step - loss: 0.3733 - accuracy: 0.8423 - val\_loss: 0.46  
42 - val\_accuracy: 0.7656  
Epoch 50/100  
44/44 [=====] - 13s 297ms/step - loss: 0.3467 - accuracy: 0.8580 - val\_loss: 0.48  
12 - val\_accuracy: 0.8062  
Epoch 51/100  
44/44 [=====] - 13s 296ms/step - loss: 0.4094 - accuracy: 0.8487 - val\_loss: 0.53  
66 - val\_accuracy: 0.7656

Epoch 52/100

44/44 [=====] - 13s 298ms/step - loss: 0.3562 - accuracy: 0.8565 - val\_loss: 0.4638 - val\_accuracy: 0.8281

Epoch 00052: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.

Epoch 53/100

44/44 [=====] - 13s 296ms/step - loss: 0.3808 - accuracy: 0.8556 - val\_loss: 0.4138 - val\_accuracy: 0.8188

Epoch 54/100

44/44 [=====] - 13s 304ms/step - loss: 0.3350 - accuracy: 0.8693 - val\_loss: 0.4298 - val\_accuracy: 0.8031

Epoch 55/100

44/44 [=====] - 13s 298ms/step - loss: 0.3095 - accuracy: 0.8778 - val\_loss: 0.4045 - val\_accuracy: 0.8500

Epoch 56/100

44/44 [=====] - 13s 302ms/step - loss: 0.3204 - accuracy: 0.8750 - val\_loss: 0.4470 - val\_accuracy: 0.8125

Epoch 57/100

44/44 [=====] - 13s 300ms/step - loss: 0.3184 - accuracy: 0.8693 - val\_loss: 0.3698 - val\_accuracy: 0.8438

Epoch 58/100

44/44 [=====] - 13s 298ms/step - loss: 0.3090 - accuracy: 0.8828 - val\_loss: 0.4346 - val\_accuracy: 0.8313

Epoch 59/100

44/44 [=====] - 13s 297ms/step - loss: 0.3140 - accuracy: 0.8800 - val\_loss: 0.3710 - val\_accuracy: 0.8531

Epoch 60/100

44/44 [=====] - 13s 298ms/step - loss: 0.2934 - accuracy: 0.8951 - val\_loss: 0.3554 - val\_accuracy: 0.8500

Epoch 61/100

44/44 [=====] - 13s 298ms/step - loss: 0.2972 - accuracy: 0.8828 - val\_loss: 0.34

40 - val\_accuracy: 0.8500  
Epoch 62/100  
44/44 [=====] - 13s 297ms/step - loss: 0.2845 - accuracy: 0.8864 - val\_loss: 0.34  
96 - val\_accuracy: 0.8594  
Epoch 63/100  
44/44 [=====] - 14s 309ms/step - loss: 0.2934 - accuracy: 0.8878 - val\_loss: 0.38  
63 - val\_accuracy: 0.8438  
Epoch 64/100  
44/44 [=====] - 13s 296ms/step - loss: 0.2947 - accuracy: 0.8800 - val\_loss: 0.36  
78 - val\_accuracy: 0.8500  
Epoch 65/100  
44/44 [=====] - 13s 299ms/step - loss: 0.2599 - accuracy: 0.8963 - val\_loss: 0.31  
34 - val\_accuracy: 0.8594  
Epoch 66/100  
44/44 [=====] - 13s 300ms/step - loss: 0.2691 - accuracy: 0.8849 - val\_loss: 0.30  
37 - val\_accuracy: 0.8938  
Epoch 67/100  
44/44 [=====] - 13s 298ms/step - loss: 0.2979 - accuracy: 0.8736 - val\_loss: 0.36  
40 - val\_accuracy: 0.8344  
Epoch 68/100  
44/44 [=====] - 13s 297ms/step - loss: 0.2769 - accuracy: 0.8913 - val\_loss: 0.36  
24 - val\_accuracy: 0.8469  
Epoch 69/100  
44/44 [=====] - 13s 295ms/step - loss: 0.2567 - accuracy: 0.8987 - val\_loss: 0.37  
48 - val\_accuracy: 0.8375  
Epoch 70/100  
44/44 [=====] - 13s 297ms/step - loss: 0.2648 - accuracy: 0.8984 - val\_loss: 0.38  
37 - val\_accuracy: 0.8531  
Epoch 71/100  
44/44 [=====] - 13s 300ms/step - loss: 0.2444 - accuracy: 0.8937 - val\_loss: 0.49  
85 - val\_accuracy: 0.8094

Epoch 72/100

44/44 [=====] - 13s 297ms/step - loss: 0.2773 - accuracy: 0.8984 - val\_loss: 0.4270 - val\_accuracy: 0.7969

Epoch 00072: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.

Epoch 73/100

44/44 [=====] - 13s 301ms/step - loss: 0.2776 - accuracy: 0.8835 - val\_loss: 0.3130 - val\_accuracy: 0.8844

Epoch 74/100

44/44 [=====] - 13s 305ms/step - loss: 0.2672 - accuracy: 0.8885 - val\_loss: 0.3974 - val\_accuracy: 0.8375

Epoch 75/100

44/44 [=====] - 13s 297ms/step - loss: 0.2516 - accuracy: 0.9091 - val\_loss: 0.3446 - val\_accuracy: 0.8438

Epoch 76/100

44/44 [=====] - 13s 303ms/step - loss: 0.2417 - accuracy: 0.9062 - val\_loss: 0.3197 - val\_accuracy: 0.8562

Epoch 77/100

44/44 [=====] - 13s 296ms/step - loss: 0.2447 - accuracy: 0.8999 - val\_loss: 0.2966 - val\_accuracy: 0.8813

Epoch 78/100

44/44 [=====] - 13s 300ms/step - loss: 0.2802 - accuracy: 0.8857 - val\_loss: 0.3149 - val\_accuracy: 0.8719

Epoch 79/100

44/44 [=====] - 13s 301ms/step - loss: 0.2409 - accuracy: 0.9027 - val\_loss: 0.3412 - val\_accuracy: 0.8625

Epoch 80/100

44/44 [=====] - 13s 297ms/step - loss: 0.2369 - accuracy: 0.9013 - val\_loss: 0.3838 - val\_accuracy: 0.8594

Epoch 81/100

44/44 [=====] - 13s 298ms/step - loss: 0.2518 - accuracy: 0.8991 - val\_loss: 0.36

63 - val\_accuracy: 0.8656

Epoch 82/100

44/44 [=====] - 13s 297ms/step - loss: 0.2408 - accuracy: 0.9034 - val\_loss: 0.30

85 - val\_accuracy: 0.8687

Epoch 83/100

44/44 [=====] - 13s 299ms/step - loss: 0.2283 - accuracy: 0.9073 - val\_loss: 0.31

11 - val\_accuracy: 0.8687

Epoch 00083: ReduceLROnPlateau reducing learning rate to 1.6000001778593287e-06.

Epoch 84/100

44/44 [=====] - 13s 297ms/step - loss: 0.2557 - accuracy: 0.9048 - val\_loss: 0.30

51 - val\_accuracy: 0.8813

Epoch 85/100

44/44 [=====] - 13s 297ms/step - loss: 0.2346 - accuracy: 0.9048 - val\_loss: 0.31

13 - val\_accuracy: 0.8656

Epoch 86/100

44/44 [=====] - 13s 296ms/step - loss: 0.2279 - accuracy: 0.9181 - val\_loss: 0.34

02 - val\_accuracy: 0.8656

Epoch 87/100

44/44 [=====] - 13s 298ms/step - loss: 0.2146 - accuracy: 0.9155 - val\_loss: 0.29

61 - val\_accuracy: 0.8625

Epoch 88/100

44/44 [=====] - 13s 295ms/step - loss: 0.2615 - accuracy: 0.9020 - val\_loss: 0.25

11 - val\_accuracy: 0.9062

Epoch 89/100

44/44 [=====] - 13s 298ms/step - loss: 0.2523 - accuracy: 0.9048 - val\_loss: 0.28

64 - val\_accuracy: 0.8875

Epoch 90/100


44/44 [=====] - 13s 296ms/step - loss: 0.2273 - accuracy: 0.9112 - val\_loss: 0.28

06 - val\_accuracy: 0.8938


Epoch 91/100

44/44 [=====] - 13s 300ms/step - loss: 0.2522 - accuracy: 0.8913 - val\_loss: 0.29  
90 - val\_accuracy: 0.8813  
Epoch 92/100  
44/44 [=====] - 13s 303ms/step - loss: 0.2485 - accuracy: 0.9009 - val\_loss: 0.30  
52 - val\_accuracy: 0.8781  
Epoch 93/100  
44/44 [=====] - 13s 297ms/step - loss: 0.2476 - accuracy: 0.9062 - val\_loss: 0.34  
75 - val\_accuracy: 0.8406  
Epoch 94/100  
44/44 [=====] - 13s 302ms/step - loss: 0.2579 - accuracy: 0.9027 - val\_loss: 0.32  
39 - val\_accuracy: 0.8625  
  
Epoch 00094: ReduceLROnPlateau reducing learning rate to 3.200000264769187e-07.  
Epoch 95/100  
44/44 [=====] - 13s 296ms/step - loss: 0.2614 - accuracy: 0.8966 - val\_loss: 0.37  
22 - val\_accuracy: 0.8469  
Epoch 96/100  
44/44 [=====] - 13s 294ms/step - loss: 0.2509 - accuracy: 0.9088 - val\_loss: 0.29  
69 - val\_accuracy: 0.8938  
Epoch 97/100  
44/44 [=====] - 13s 297ms/step - loss: 0.2205 - accuracy: 0.9169 - val\_loss: 0.25  
75 - val\_accuracy: 0.9000

Training Loss VS Validation Loss



Training Accuracy VS Validation Accuracy





Model saved successfully!


Confusion matrix for train data:

Found 5712 images belonging to 4 classes.

179/179 [=====] - 23s 126ms/step - loss: 0.1976 - accuracy: 0.9321

Score = [0.1975684016942978, 0.9320728182792664]

Accuracy = 0.9320728291316527



Confusion matrix for val data:

Found 1311 images belonging to 4 classes.

41/41 [=====] - 5s 111ms/step - loss: 0.3100 - accuracy: 0.8757

Score = [0.31003186106681824, 0.8756674528121948]

Accuracy = 0.8756674294431731

